



Multivalue Fields

Document Usage Guidelines

- Should be used only for enrolled students
- Not meant to be a self-paced document, an instructor is needed
- Lab Exercise slides reference the hands-on lab exercise guide
- Do not distribute

Course Goals

- Define multivalue fields
- Evaluate multivalue fields
- Manipulate multivalue fields
- Analyze multivalue fields
- Format multivalue fields
- Configure extractions of multivalue fields

Course Outline

- What are Multivalue Fields?
- Create Multivalue Fields
- Evaluate Multivalue Fields
- Analyze Multivalue Data

What are Multivalue Fields?

Topic Objectives

- Define multivalue fields
- Define self-describing data
- Understand how JSON data is handled in Splunk
- Use the **spath** command to interpret self-describing data
- Manipulate multivalue fields with **mvzip** and **mvexpand**
- Convert single-value fields to multivalue fields with:
 - **makemv** command
 - **stats list** and **values** functions
 - **transaction** command

Define Multivalue Fields

- Multivalue fields are fields that hold more than one value:
 - Email logs containing To, Cc, and Bcc values
 - AWS configuration records
 - JSON and XML arrays
- Multivalue functions and commands provide ways to count, evaluate, modify, create, combine, sort, and split multivalue fields
 - Typically contain a `mv` prefix
 - Exceptions discussed in this module:
 - `split`, `list`, and `values` functions
 - `transaction` and `multikv` commands

Define Self-Describing Data

- JSON and XML data are types of "self-describing data" where the structure is embedded in the data itself
- Comprised of metadata which may include:
 - Properties/elements/attributes
 - Data types/items
 - Compression/encoding scheme
 - Other info

JSON and XML Examples

JSON

```
{
  "menu": {
    "id": "file",
    "value": "File",
    "popup": {
      "menuitem": [
        {
          "value": "New", "onclick": "CreateNewDoc()"
        },
        {
          "value": "Open", "onclick": "OpenDoc()"
        },
        {
          "value": "Close", "onclick": "CloseDoc()"
        }
      ]
    }
  }
}
```

XML

```
<menu id="file" value="File">
  <popup>
    <menuitem value="New" onclick="CreateNewDoc()" />
    <menuitem value="Open" onclick="OpenDoc()" />
    <menuitem value="Close" onclick="CloseDoc()" />
  </popup>
</menu>
```

- Splunk extracts fields from JSON and XML based on the formatting
- Both consist of:
 - Collections of name/value pairs
 - Ordered lists of values (arrays)

Splunk and JSON Data

- Splunk preserves JSON structure when displaying data from JSON-formatted log files
- In JSON data:
 - { } indicates an object, which is a grouping of field-value pairs
 - [] indicates an array of objects
- When extracted by Splunk, the array contents become multivalue fields

index=systems sourcetype=system_info

i	Time	Event
>	6/21/22 1:59:10.000 PM	{ [-] CPU: { [+] } CPU_CORES: [[+]] RAM: { [+] } ROOT_USERS: [[+]] SYSTEM: [[+]] asctime: 2022-06-21T19:59:10+00:00 created: 1655841550.118906 levelname: INFO message: system_info name: SystemLogger

JSON event data

Automatically Extracting JSON Data

index=systems sourcetype=system_info

i	Time	Event
>	6/21/22 1:59:10.000 PM	<pre>{ [-] CPU: { [+] } CPU_CORES: [[-] { [-] core: core_number_1 core_percent_used: 0 guest: 0 guest_nice: 0 idle: 100 iowait: 0 irq: 0 nice: 0 softirq: 0 steal: 0 system: 0 user: 0 } RAM: { [+] } ROOT_USERS: [[-] Admin Duke Zed Camilian] SYSTEM: [[+] }</pre>

The object array **CPU_CORES[]** has multiple objects with the same fields. When *extracted*, these fields are represented as multivalue fields.

When *ingested*, the array **ROOT_USERS[]** becomes a multivalue field with 4 values

```
ROOT_USERS: [ [-]
Duke
Camilian
Admin
Cardozo
```

```
CPU_CORES: [ [-]
{ [-]
  core: core_number_1
  core_percent_used: 0
  guest: 0
  guest_nice: 0
  idle: 100
  iowait: 0
  irq: 0
  nice: 0
  softirq: 0
  steal: 0
  system: 0
  user: 0
}
```

```
a CPU_CORES{}.core 16
a CPU_CORES{}.core_percent_used 11
a CPU_CORES{}.guest 1
a CPU_CORES{}.guest_nice 1
a CPU_CORES{}.idle 11
a CPU_CORES{}.iowait 2
a CPU_CORES{}.irq 1
a CPU_CORES{}.nice 3
a CPU_CORES{}.softirq 2
a CPU_CORES{}.steal 2
a CPU_CORES{}.system 4
a CPU_CORES{}.user 8
```

Interesting Fields List

JSON Field Name with {} Caveat

- Functions and commands process {} within a field name as special characters
- For consistent results, use the **rename** command to remove the {} entirely from the field name

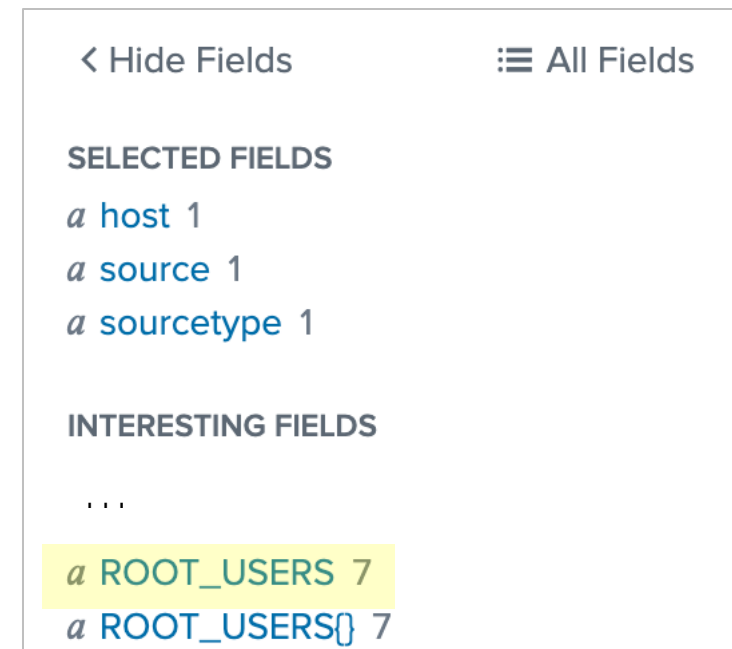
```
index=systems sourcetype=system_info  
| rename ROOT_USERS{} as root_users  
| eval root_users=mvsort(root_users)
```

JSON Field Name with {} Caveat (cont.)

- Using `rename` is recommended because certain commands produce inconsistent results when used with `{}`-containing fields
- For example, the `eval` statement *seems* to overwrite `ROOT_USERS{}`

```
index=systems sourcetype=system_info  
| eval ROOT_USERS{}=upper('ROOT_USERS{'})
```

- Instead, it creates a new field, `ROOT_USERS` (without the `{}`)



Interpreting XML

By default, Splunk doesn't extract fields automatically from XML data

< Hide Fields		≡ All Fields	i	Time	Event
SELECTED FIELDS			>	6/21/22 1:40:53.000 PM	<?xml version="1.0" encoding="UTF-8"?> <root> <row> <status>100</status> <status_description>Continue</status_description> <status_type>Informational</status_type> </row> <row> <status>101</status> <status_description>Switching Protocols</status_description> <status_type>Informational</status_type> </row> <row> <status>200</status>
INTERESTING FIELDS					
a encoding 1					
a index 1					
# linecount 1					
a punct 1					
a splunk_server 1					
a timestamp 1					
# version 1					

Note

Admins can configure the sourcetype with `KV_MODE=XML` to automatically extract fields at search time.

Use `spath` with XML

The `spath` command interprets the XML structure so you can access the data as Splunk fields

```
index=systems sourcetype=status_definitions  
| spath
```

<div><div>< Hide Fields</div><div>☰ All Fields</div></div>	<div><div><div>i</div><div>Time</div><div>Event</div></div></div>	
<div><div>SELECTED FIELDS</div><div><div>a host 1</div><div>a source 1</div><div>a sourcetype 1</div></div></div> <div><div>INTERESTING FIELDS</div><div><div>a encoding 1</div><div>a index 1</div><div># linecount 1</div><div>a punct 1</div><div><div>a root.row.status 40</div><div>a root.row.status_description 40</div><div>a root.row.status_type 5</div></div><div>a splunk_server 1</div><div>a timestamp 1</div><div># version 1</div></div></div>	<div><div>▼</div><div>6/21/22 1:40:53.000 PM</div><div><?xml version="1.0" encoding="UTF-8"?> <root> <row> <status>100</status> <status_description>Continue</status_description> <status_type>Informational</status_type> </row> <row> <status>101</status> <status_description>Switching Protocols</status_description> <status_type>Informational</status_type> </row> <row> <status>200</status> <status_description>OK</status_description> <status_type>Successful</status_type> </row></div></div>	<div><div>Note</div><div>Field name parent ele</div></div>

spath Command

```
... | spath [input=<field>][output=<field>][path=<datapath>|<datapath>]
```

- Extracts fields from self-describing data (XML and JSON)
- Can be used alone or with optional arguments:
 - **input**: specifies which **<field>** to extract data from (defaults to **_raw**)
 - **output**: data to be extracted is written to this **<field>** (defaults to value of the **path** argument)
 - **path**: the location path to the value you want to extract
 - Valid syntax is **path=<datapath>** or just **<datapath>**
 - By default, extracts all fields from first 5000 characters of **input**

spath Command: Location Steps

- The **path** argument can contain one or more location steps, separated by periods
- Each step contains field name and optional index (position) in curly brackets
 - If index is an integer, it specifies the position of data in an array
 - If index is a string preceded by an **@** symbol, it specifies an XML attribute
- Examples:
 - `recordings.album.artist`
 - `entities.hashtags{3}.text`
 - `purchases.book.title{@yearPublished}`

Note



Indexes in JSON and XML are slightly different. In JSON, numbering begins with 0; in XML, it begins with 1.

spath Example 1

Scenario



IT wants to create a table containing status, description, and status type from data in an XML file.

```
index=systems sourcetype=status_definitions
| spath
| table root.*
| rename root.row.* as *
```

status	status_description	status_type
100	Continue	Informational
101	Switching Protocols	Informational
200	OK	Successful
201	Created	Successful
202	Accepted	Successful
203	Non-Authoritative Information	Successful
204	No Content	Successful
205	Reset Content	Successful
206	Partial Content	Successful
300	Multiple Choices	Redirection
301	Moved Permanently	Redirection
302	Found	Redirection
303	See Other	Redirection
304	Not Modified	Redirection

spath Example 2

Use curly brackets without an integer to indicate the index of an entire array

Scenario



The Dev team wants to display a table of popup menu values from data in a JSON file.

Note



If you do not use the **output** option, then you should rename the field to remove the {}. Doing so avoids issues that can occur when using field names containing {} with certain commands.

```
index=systems sourcetype=menu_json  
| spath output=menuItems path=menu.popup.menuitem{}.value  
| table menuItems
```

menuItems ▾

New

Open

Close

```
{  
  "menu": {  
    "id": "file",  
    "value": "File",  
    "popup": {  
      "menuitem": [  
        {"value": "New", "onclick": "CreateNewDoc()"},  
        {"value": "Open", "onclick": "OpenDoc()"},  
        {"value": "Close", "onclick": "CloseDoc()"}  
      ]  
    }  
  }  
}
```

spath Example 3

Use the @ symbol to specify an XML attribute

Scenario



The Docs team wants to extract a table of book publication dates from an XML file.

```
index=systems sourcetype=library_xml
| spath output=publicationDates path=purchases.book.title{@yearPublished}
| table publicationDates
```

```
<purchases>
  <book>
    <title yearPublished=1867>War and Peace</title>
    <author>Leo Tolstoy</author>
  </book>
  <book>
    <title yearPublished=1866>Crime and Punishment</title>
    <author>Fyodor Dostoyevsky</author>
  </book>
  <book>
    <title yearPublished=1877>Anna Karenina</title>
    <author>James Joyce</author>
  </book>
  <book>
    <title yearPublished=1880>The Brothers Karamazov</title>
    <author>Fyodor Dostoyevsky</author>
  </book>
  <book>
    <title yearPublished=1851>Moby Dick</title>
    <author>Herman Melville</author>
  </book>
</purchases>
```

publicationDates

1867
1866
1877
1880
1851

spath Example 4

Even within unstructured data, a field can have a valid self-describing data structure that can be extracted using **spath**

The screenshot shows the Splunk search results interface. On the left, the 'SELECTED FIELDS' list includes 'a system_info 100+'. The main panel displays the 'system_info' field with a 'Top 10 Values' table. A callout box highlights the search path `index=systems sourcetype=server_log`.

Top 10 Values	Count	%
<code>{"CPU Percent Used": 0.0, "RAM Percent Used": 11.7, "Free Ram": 2507464704, "Used Ram": 211558400, "Available Ram": 3644100608, "Active Ram": 612716544, "Inactive Ram": 782462976, "Total Ram": 4124807168}</code>	201	0.077%
<code>{"CPU Percent Used": 100.0, "RAM Percent Used": 12.8, "Free Ram": 1786740736, "Used Ram": 223539200, "Available Ram": 3597025280, "Active Ram": 857894912, "Inactive Ram": 1179504640, "Total Ram": 4124807168}</code>	198	0.076%
<code>{"CPU Percent Used": 0.0, "RAM Percent Used": 11.6, "Free Ram": 2507730944, "Used Ram": 211337216, "Available Ram": 3644321792, "Active Ram": 612634624, "Inactive Ram": 782438400, "Total Ram": 4124807168}</code>	134	0.052%
<code>{"CPU Percent Used": 0.0, "RAM Percent Used": 11.6, "Free Ram": 2510512128, "Used Ram":</code>	134	0.052%

The `system_info` field has a valid JSON structure

spath Example 4 (cont.)

To extract data from the `system_info` field, invoke the `spath` command with the `input` option

< Hide Fields

≡ All Fields

SELECTED FIELDS

a host 1

a source 1

a sourcetype 1

INTERESTING FIELDS

Active Ram 100+

Available Ram 100+

CPU Percent Used 100+

date_hour 24

date_mday 31

date_minute 60

a date_month 2

date_second 24

a date_wday 7

date_year 1

date_zone 1

Free Ram 100+

Inactive Ram 100+

a index 1

linecount 1

a punct 1

RAM Percent Used 18

a splunk_server 1

a system_info 100+

a system_name 1

timeendpos 1

timestartpos 1

Total Ram 1

Used Ram 100+

i	Time	Event
>	6/21/22 2:23:50.000 PM	[2022-06-21T20:23:50+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 13.9, "RAM Percent Used": 11.7, "Free Ram": 2504822784, "Used Ram": 211955712, "Available Ram": 3643686912, "Active Ram": 613863424, "Inactive Ram": 783532032, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log
>	6/21/22 2:23:40.000 PM	[2022-06-21T20:23:40+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 16.9, "RAM Percent Used": 11.7, "Free Ram": 2504822784, "Used Ram": 211955712, "Available Ram": 3643686912, "Active Ram": 613863424, "Inactive Ram": 783532032, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log
>	6/21/22 2:23:30.000 PM	[2022-06-21T20:23:30+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 10.0, "RAM Percent Used": 11.7, "Free Ram": 2504822784, "Used Ram": 211968000, "Available Ram": 3643686912, "Active Ram": 613863424, "Inactive Ram": 783532032, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log
>	6/21/22 2:23:20.000 PM	[2022-06-21T20:23:20+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 17.9, "RAM Percent Used": 11.7, "Free Ram": 2505338880, "Used Ram": 211456000, "Available Ram": 3644198912, "Active Ram": 613863424, "Inactive Ram": 783527936, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log
>	6/21/22 2:23:10.000 PM	[2022-06-21T20:23:10+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 18.0, "RAM Percent Used": 11.7, "Free Ram": 2505338880, "Used Ram": 211468288, "Available Ram": 3644194816, "Active Ram": 613863424, "Inactive Ram": 783523840, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log
>	6/21/22 2:23:00.000 PM	[2022-06-21T20:23:00+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 6.5, "RAM Percent Used": 11.7, "Free Ram": 2505338880, "Used Ram": 211468288, "Available Ram": 3644190720, "Active Ram": 613826560, "Inactive Ram": 783515648, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log
>	6/21/22 2:22:50.000 PM	[2022-06-21T20:22:50+00:00] INFO [muddledLogger] system_name="Web Server", system_info={"CPU Percent Used": 18.0, "RAM Percent Used": 11.7, "Free Ram": 2505338880, "Used Ram": 211472384, "Available Ram": 3644190720, "Active Ram": 613826560, "Inactive Ram": 783515648, "Total Ram": 4124807168} host = mixed_system_data source = /opt/log/muddled.log sourcetype = server_log

index=systems sourcetype=server_log
| **spath input=system_info**

Use the eval Command spath Function

- As an alternative to the **spath** command, you can also use the **spath** function of the **eval** command

```
index=systems sourcetype=status_definitions  
| eval description = upper(spath(_raw, "root.row.status_description"))  
| table description
```

- spath(X,Y)** where:
X: input source field
Y: XML or JSON
formatted location path
to the value you want to
extract from X

description ▾


- CONTINUE
- SWITCHING PROTOCOLS
- OK
- CREATED
- ACCEPTED
- NON-AUTHORITATIVE INFORMATION
- NO CONTENT
- RESET CONTENT
- PARTIAL CONTENT
- MULTIPLE CHOICES
- MOVED PERMANENTLY

Extract Fields Using the multikv Command

- For table-formatted events, **multikv** creates an event for each row
- Field names derived from header row of each event

INTERESTING FIELDS

- `a eventtype 1`
- `a ForeignAddress 46`
- `a index 1`
- `# linecount 1`
- `a LocalAddress 100+`
- `a Proto 2`
- `a punct 1`
- `# Recv_Q 1`
- `# Send_Q 8`
- `a splunk_server 1`
- `a State 4`
- `a tag 5`
- `a tag::eventtype 5`
- `a timestamp 1`



i	Time	Event						
>	8/8/18 10:26:52.000 AM	Proto	Recv-Q	Send-Q	LocalAddress	ForeignAddress	State	
		udp	0	0	*:68	*:*	<n/a>	
		udp	0	0	*:111	*:*	<n/a>	
		udp	0	0	10.0.1.80:123	*:*	<n/a>	
		udp	0	0	127.0.0.1:123	*:*	<n/a>	
		Show all 54 lines						
		host = sh-20711 source = netstat sourcetype = netstat						

index=main sourcetype=netstat

>	8/8/18 10:41:52.000 AM	udp	0	0	*:68	*:*		
		host = sh-20711 source = netstat sourcetype = netstat						
>	8/8/18 10:41:52.000 AM	udp	0	0	*:111	*:*	<n/a>	
		host = sh-20711 source = netstat sourcetype = netstat						
>	8/8/18 10:41:52.000 AM	udp	0	0	10.0.1.80:123	*:*	<n/a>	
		host = sh-20711 source = netstat sourcetype = netstat						
>	8/8/18 10:41:52.000 AM	udp	0	0	127.0.0.1:123	*:*	<n/a>	
		host = sh-20711 source = netstat sourcetype = netstat						

index=main sourcetype=netstat
| multikv

Use multikv: fields and filter Options

- Use `fields` option to extract only specified fields
- Use `filter` option to include only table rows containing at least one field value from a specified list of fields

```
index=main sourcetype=netstat
| multikv fields LocalAddress ForeignAddress State filter ESTAB LISTEN
| table LocalAddress, ForeignAddress, State
```

LocalAddress	ForeignAddress	State
*:46449	*:*	LISTEN
*:22	*:*	LISTEN
127.0.0.1:8191	127.0.0.1:38404	ESTAB
127.0.0.1:38404	127.0.0.1:8191	ESTAB
127.0.0.1:38430	127.0.0.1:8191	ESTAB

Note

The arguments supplied to the `filter` option are treated like they are connected using OR operators.

Multivalue Function: **list**

```
...| stats list(X) [as <newfield>| by <field>]
```

- A multivalue function that can be used with **stats** or **chart**
- Lists all field values for a given field, **X**
 - Listed fields can be renamed as **<newfield>**, otherwise the field name will be **list(X)**
 - Values can be grouped by **<field>**
 - If more than one value is listed, the result is a multivalue field
- The order of the values reflects the order of events

Multivalue Function: list Example

Scenario



Which websites has each employee accessed during the last 60 minutes?

```
index=network sourcetype=cisco_wsa_squid  
| stats list(s_hostname) as "Websites visited:"  
by cs_username
```

cs_username	Websites visited:
acurry@buttercupgames.com	- - www.goppo.com www.goppo.com
bgenin@buttercupgames.com	www.hybridarcade.com
cberztiss@buttercupgames.com	- -
cganttchart@buttercupgames.com	-
djohnson@buttercupgames.com	-
dpiazza@buttercupgames.com	-
fullian@buttercupgames.com	- www.espn.go.com www.espn.go.com
gbottazzi@buttercupgames.com	- -
gfacello@buttercupgames.com	www.fftoday.com www.fftoday.com www.fftoday.com www.fftoday.com www.fftoday.com

Multivalue field

Multivalue Function: `values`

```
...| stats values(X) [by <field> | as <newfield>]
```

- A multivalue function that can be used with `stats` or `chart`
- Syntax is identical to `list` function
- Unlike the `list` function, the `values` function lists all *unique* values of field, `X`

Multivalue Functions: list vs values

```
index=network sourcetype=cisco_wsa_squid  
| stats list(s_hostname) as "Websites visited:"  
by cs_username
```

cs_username	Websites visited:
adombrowski@buttercupgames.com	www.ebgames.com www.ebgames.com
apucci@buttercupgames.com	www.ebgames.com www.ebgames.com www.ebgames.com
blu@buttercupgames.com	www.ebgames.com
cberztiss@buttercupgames.com	www.ebgames.com
cmunson@buttercupgames.com	www.collegegrad.com
cquinn@buttercupgames.com	www.ebgames.com www.ebgames.com
djohnson@buttercupgames.com	www.ebgames.com
edutra@buttercupgames.com	www.ebgames.com
ewarwick@buttercupgames.com	www.collegegrad.com www.ebgames.com

```
index=network sourcetype=cisco_wsa_squid  
| stats values(s_hostname) as "Websites visited:"  
by cs_username
```

cs_username	Websites visited:
adombrowski@buttercupgames.com	www.ebgames.com
apucci@buttercupgames.com	www.ebgames.com
blu@buttercupgames.com	www.ebgames.com
cberztiss@buttercupgames.com	www.ebgames.com
cmunson@buttercupgames.com	www.collegegrad.com
cquinn@buttercupgames.com	www.ebgames.com
djohnson@buttercupgames.com	www.ebgames.com
edutra@buttercupgames.com	www.ebgames.com
ewarwick@buttercupgames.com	www.collegegrad.com www.ebgames.com

transaction Command

When processed by the `transaction` command, single-value fields become multivalue fields (if there are several values)

```
index=web sourcetype=access*
| transaction clientip
| table clientip status action
```

clientip	status	action
221.207.229.6	200	addtocart view
118.142.68.222	200 400	addtocart purchase view
12.130.60.4	200	changequantity view
99.61.68.230	200	addtocart changequantity purchase view
58.68.236.98	200	addtocart purchase

What are Multivalue Fields Lab Exercise

Time: 20 minutes

Tasks:

- Extract fields from an XML file using the **spath** command and the **spath** function of the **eval** command
- Use the **spath** command to extract fields from the **system_info** field and use this data to display server performance
- Complete a search with a multivalue **stats** function
- Use a multivalue **stats** function to list all unique users active on the AD/DNS server during the last 4 hours

Create Multivalue Fields

Topic Objectives

- Create multivalue fields with:
 - `split` (multivalue `eval` function)
 - `makemv` (multivalue command)

Multivalue eval Functions

The `eval` functions discussed in this topic (and succeeding topics) can all be used with `eval`, `fieldformat`, and `where` commands...

```
...| eval <field>=mvfunction(...)
```

```
...| where <booleanExpression>
```

The `booleanExpression` would contain a `mvfunction`

```
...| fieldformat <field>=mvfunction(...)
```

...and as part of `eval` expressions

```
...| stats function(eval(mvfunction(...)))
```

Create Multivalue Fields with `split`

```
...| eval <field>=split(<mvfield>,"<delimiter>")
```

- Multivalue function
- Separates a single-value field based on a given delimiter and returns a multivalue field as result
 - <mvfield> is a single-value field
 - "<delimiter>" is the delimiter that <mvfield> is split with

```
...| eval newfield=split(fieldA,";")  
| table fieldA newfield
```

fieldA	newfield
value1;value2;value3	value1 value2 value3

split Example

Scenario



For the **Last 7 days**, Sales wants to extract each part of the **productId** as a separate value within a multivalue field in order to capture a new field called **productCodes** (The **productId** has the format AA-BB-CCC.)

```
index=web sourcetype=access_combined
action=purchase status=200 productId!=NULL
| eval productCodes = split(productId, "-")
| table productId, productCodes, product_name
```

productId	productCodes	product_name
FS-SG-G03	FS SG G03	Final Sequel
MB-AG-T01	MB AG T01	Manganiello Bros. Tee
DB-SG-G01	DB SG G01	Mediocre Kingdoms

Convert Single Value Fields with makemv

```
... | makemv [delim=<string> |tokenizer=<regex>] <field>
```

- Multivalue command
- Converts an *existing* single value field to a multivalue field, based on a provided string or regex
 - `delim` specifies a string delimiter; defaults to a single space “ “
 - `tokenizer` specifies a regular expression
- `makemv` is a distributable streaming command

makemv Example

makemv captures words and makes them into values of mv_uri_path

```
index=web uri_path=*  
| table uri_path  
| dedup uri_path  
| eval mv_uri_path = uri_path  
| makemv tokenizer="(\w+)" mv_uri_path
```

uri_path	mv_uri_path
/cart/success.do	cart success do
/cart.do	cart do
/productscreen.html	productscreen html
/cart/error.do	cart error do

Evaluate Multivalue Fields

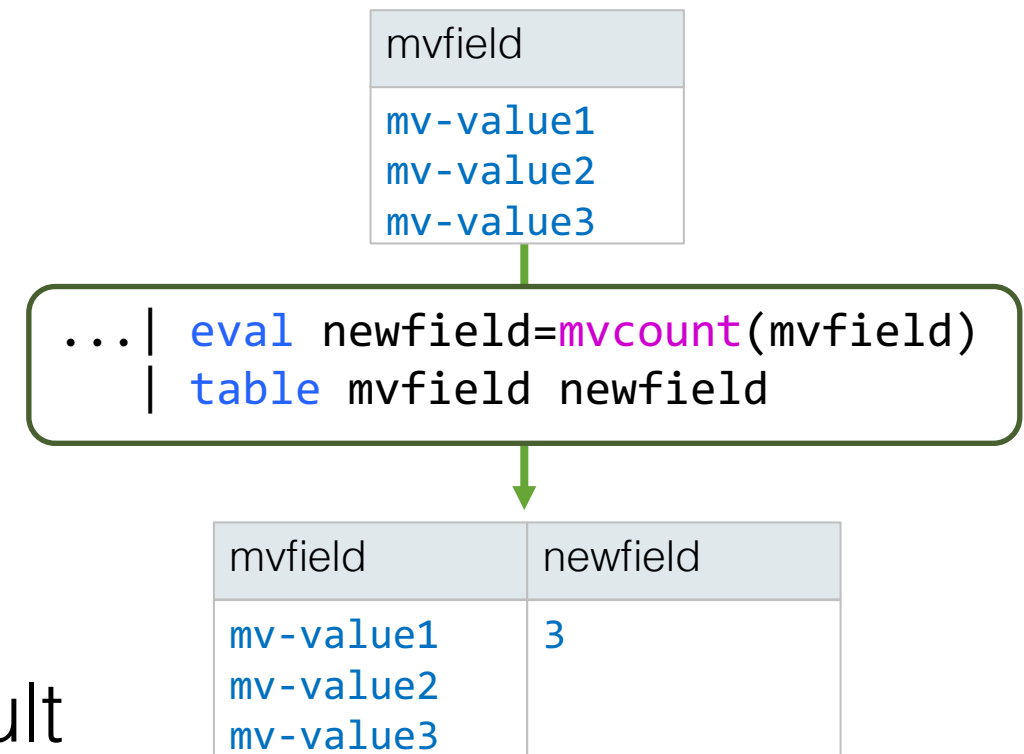
Topic Objectives

- Evaluate multivalue fields using multivalue `eval` functions:
 - `mvcount`
 - `mvindex`
 - `mvfilter`

Count Values with mvcount

```
...| eval <field>=mvcount(<mvfield>)
```

- Multivalue function
- Counts the number of values for a specified multivalue field, <mvfield>
- Returns NULL when:
 - The field has no values
 - If a field does not exist for a particular result



mvcount Example

Scenario



Sales wants to know how many transactions in the online store over the last 24 hours contained more than 4 actions within half an hour.

```
index=web sourcetype=access_combined
| transaction JSESSIONID maxspan=30m
| eval actions=mvcount(action)
| where actions > 4
| stats count
```

count
15

action	actions
addtocart	2
purchase	
addtocart	2
purchase	
addtocart	3
purchase	
view	

This is what you would see if you put **action** and **actions** in a table

Pull Specific Values Using `mvindex`

```
... | eval <field>=mvindex(<mvfield>,startIndex,[endIndex])
```

- Multivalue function
- Assigns an array index to a multivalue field and returns a value based on a given integer
 - `<mvfield>` is a multivalue field
 - `startIndex` and `endIndex` are integers referencing the position of a value within a multivalue field index
 - If `endIndex` is specified, Splunk returns values between the `startIndex` integer and `endIndex` integer
 - If `endIndex` is not specified, only the value at `startIndex` is returned

productCode	
PZ	index=0
SG	index=1
G05	index=2

mvindex Example

Scenario



Sales wants a count of all product names (product_name) sold yesterday on the Buttercup Games website whose second productCodes value is SH. (Recall that productCodes is created from productId.)

```
index=web sourcetype=access_combined action=purchase
status=200 productId!=NULL
| eval productCodes = split(productId, "-")
| eval productCode2 = mvindex(productCodes, 1)
| where ( productCode2 = "SH" )
| stats count(productCode2) as count by product_name
```

This is what you would see if you put productCodes and productCode2 in a table

productCodes	productCode2
PZ SG G05	SG
DB SG G01	SG
MB AG T01	AG
FI AG G08	AG
SC MG G10	MG

product_name	count
Fire Resistance Suit of Provolone	128
Holy Blade of Gouda	92
World of Cheese	152
World of Cheese Tee	104

Filter Multivalue Field Values with `mvfilter`

```
...| eval <field>=mvfilter(<booleanExpression>)
```

- Multivalue function
- Filters a multivalue field based on an arbitrary Boolean expression
- `<booleanExpression>` can reference only one field at a time

mvfilter Example: Two Solutions

Scenario



IT wants a list of all **ROOT_USERS** whose name begins with C.

```
index=systems sourcetype=system_info
| rename ROOT_USERS{} as root_users
| eval names=mvfilter(match(root_users, "^C"))
| dedup names
| table names
```

match finds values in **root_users** that match the regex: **^C**

Use rename to make **ROOT_USERS{}** easier to use

```
index=systems sourcetype=system_info
| rename ROOT_USERS{} as root_users
| eval names=mvfilter(root_users LIKE "C%")
| dedup names
| table names
```

The **LIKE** Boolean finds values in **root_users** that start with **C**

names

Camilian

Camilian
Cardozo

Create and Evaluate Multivalue Fields Lab Exercise

Time: 15 minutes

Tasks:

- Use the `makemv` command to convert the `productId` field into a multivalue field and find all products sold yesterday whose `productId` contains "SH"
- Use various multivalue `eval` functions to complete a search that will find employees who used a workstation other than their own
- Challenge: Test the previous search by using an `eval` function to simulate the scenario of an employee logging into another employee's workstation

Analyze Multivalue Data

Topic Objectives

- Analyze multivalue data using multivalue `eval` functions:
 - `mvsort`
 - `mvzip`
 - `mvjoin`
 - `mvmap`
 - `mvappend`
- Analyze multivalue data using the `mvexpand` multivalue command

Sort Multivalue Field Values with `mvsort`

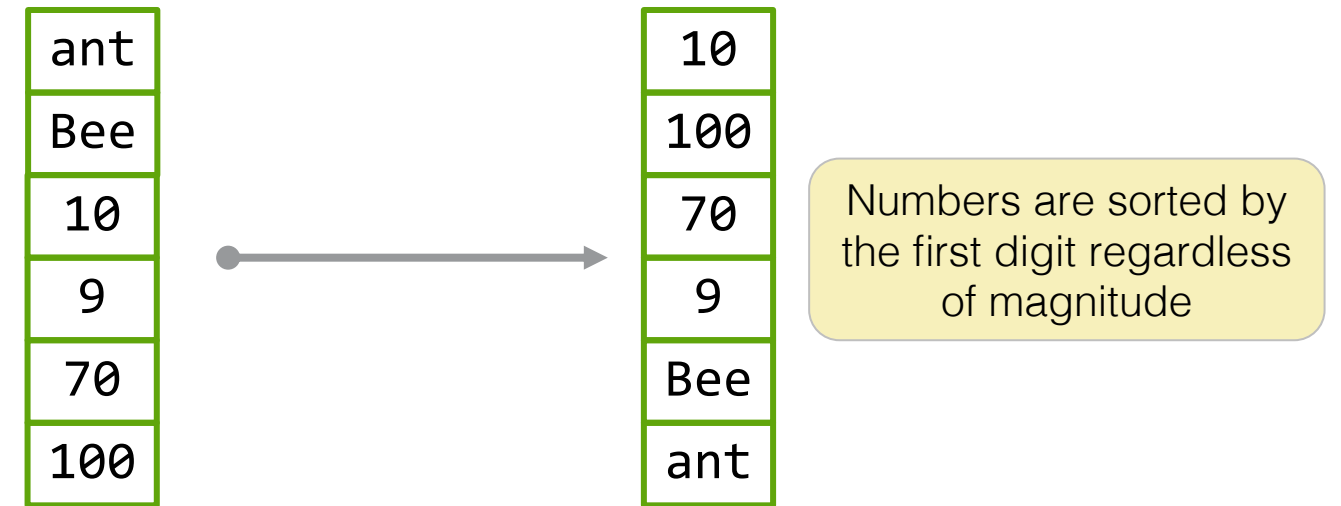
```
...| eval <field>=mvsort(<mvfield>)
```

- Multivalue function
- Returns a multivalue field with values sorted lexicographically
- `<mvfield>` is a multivalue field

mvsort Function: Lexicographical Order

Items are sorted based on their encoding; in Splunk, encoding is almost always UTF-8

- Numbers are sorted based on the first digit
- Numbers are before letters
- Uppercase letters are before lowercase letters



Note

The **mvsort** function always sorts lexicographically. This is different than the **sort** command which first determines what type of field it is sorting and then sorts based on field type.

mvsort Example

Scenario



IT wants a list of root users for each of its virtual AWS instances (i.e., systems) for the last 7 days.


```
index=systems sourcetype=system_info
| rename ROOT_USERS{} as root_users,
  SYSTEM{} as system
| eval root_users=mvsort(root_users)
| table system, root_users
| stats values(root_users) as root_users by system
```

system	root_users
nix_system_idx_1	Admin Camilian Cardozo Duke
nix_system_idx_1_UK	Admin Camilian Cardozo Duke
nix_system_idx_2	Beulah Camilian Duke Zed

Combine Multivalue Fields with mvzip

```
...| eval <field>=mvzip(X,Y["<delimiter>"])
```

- Multivalue function
- Concatenates the values of two multivalue fields
 - X and Y are multivalue fields
 - The first value of X is concatenated with the first value of Y, etc.
 - "<delimiter>" is used as the separator; if none specified, the default delimiter is a comma (,)
- Fields X and Y should have the same number of values
- Resulting field values are treated as strings



mvzip Example 1

Scenario

IT wants a report for the last 24 hours of all its AWS systems, along with each of its cores and the percent used for each core. Present the `core` and `core_percent_used` in a single field, concatenated by a colon and a space. Show only cores with usage more than 0.0.



```
index=systems sourcetype=system_info
| rename SYSTEM{} as system, CPU_CORES{}.core as core,
CPU_CORES{}.core_percent_used as percent_used
| eval zip_percent_used = mvzip(core, percent_used, ": ")
| stats count as sum_core by system, asctime,
zip_percent_used
| search zip_percent_used!="*: 0.0"
| stats list(zip_percent_used) as "CPU Core Usage" by
system, asctime
```

core	percent_used
core_number_5	1.0

system	asctime	CPU Core Usage
nix_system_idx_1	2020-09-16 23:05:16,950	core_number_5: 1.0
nix_system_idx_1	2020-09-16 23:08:05,156	core_number_2: 1.0
nix_system_idx_1	2020-09-16 23:09:01,225	core_number_8: 1.0
nix_system_idx_1	2020-09-16 23:15:33,685	core_number_6: 1.0
nix_system_idx_1	2020-09-16 23:20:14,031	core_number_12: 1.0 core_number_15: 1.0 core_number_9: 1.0

mvzip Example 2

If the X and Y fields do not have the same number of values, the results will not be complete

```
index=systems sourcetype=system_info
| rename SYSTEM{} as system
| rename ROOT_USERS{} as users
| eval systemUsers=mvzip(system,users,"-")
| table system, users, systemUsers
```

The field with fewer values determines how many values the new field will have

system	users	systemUsers
nix_system_idx_4	Admin Duke Zed Camilian	nix_system_idx_4-Admin
nix_system_idx_4_UK	Admin Duke Zed Camilian	nix_system_idx_4_UK-Admin
nix_system_idx_3	Beulah Camilian Admin Nells	nix_system_idx_3-Beulah

Concatenate with mvjoin

```
...| eval <field>=mvjoin(<mvfield>,"<delimiter>")
```

- Multivalue function
- Concatenates, i.e., links together, values of a multivalue field with a specified delimiter, which outputs to a single value field
 - <mvfield> is a multivalue field
 - "<delimiter>" is used as the separator

mvjoin Example

Scenario



IT wants to see a list of root users (concatenated with a comma) on each of its AWS instances during the last 60 minutes.

```
index=systems sourcetype=system_info
| rename ROOT_USERS{} as users, SYSTEM{} as system
| eval users=mvsort(users), users=mvjoin(users,",")
| stats values(system) as systems by users
```

users ▾	systems ▾
Admin,Beulah,Camilian,Nells	nix_system_idx_3
Admin,Camilian,Cardozo,Duke	nix_system_idx_1
Admin,Camilian,Duke,Zed	nix_system_idx_4
Beulah,Camilian,Duke,Zed	nix_system_idx_2

Operate on Multivalue Fields with mvmap

```
...| eval <evalField>=mvmap(<mvfield>,<expression>)
```

- Multivalue function
- Performs an operation over the values of a multivalue field then returns a multivalue field as a result
 - <mvfield> is the multivalue field to be operated on
 - <expression> provides the operation to execute

mvmmap Example

```
index=systems sourcetype=system_info
| rename SYSTEM{} as system, CPU_CORES{}.core as core,
  CPU_CORES{}.core_percent_used as percent_used
| eval percent_used = mvmmap(percent_used, percent_used/2)
| eval zip_percent_used = mvzip(core, percent_used, ": ")
| stats count as sum_core by system, asctime,
  zip_percent_used
| search zip_percent_used!="*: 0.0"
| stats list(zip_percent_used) as "CPU Core Usage" by
  system, asctime
```

mvmmap applies the operation (dividing by 2) to all the values of **percent_used**

system	asctime	CPU Core Usage
nix_system_idx_1	2019-09-06 17:45:26,956	core_number_11: 0.50
nix_system_idx_1	2019-09-06 17:51:03,371	core_number_4: 0.50
nix_system_idx_1	2019-09-06 17:53:51,578	core_number_11: 0.50
nix_system_idx_1	2019-09-06 18:02:16,201	core_number_9: 0.50

Combine Fields and Values with mvappend

```
...| eval <field>=mvappend(arg1, arg2, arg3, ..., argN)
```

- Multivalue function
- Takes two or more single-value fields, multivalue fields, and strings and returns a multivalue field as the result
- Each input becomes a separate value within the new multivalue field

fieldA	mvfield
value	mv-value1 mv-value2

```
...| eval newfield=mvappend(fieldA,"abc",mvfield)  
| table fieldA mvfield newfield
```

fieldA	mvfield	newfield
value	mv-value1 mv-value2	value abc mv-value1 mv-value2

mvappend Example

```
| makeresults
| eval manufacturer = "Donco",
    city = "Shenzen",
    country = "China",
    products = mvappend("Holy Blade of Gouda", "Mediocre Kingdoms",
        "Dream Crusher")
```

_time ▾	city ▾	country ▾	manufacturer ▾	products ▾
2022-06-21 15:17:26	Shenzen	China	Donco	Holy Blade of Gouda Mediocre Kingdoms Dream Crusher

Note

`makeresults` is a useful command for creating data to test and troubleshoot with.

The new values of the multivalue field, **products**

Analyze Multivalue Fields with mvexpand

```
...| mvexpand <mvfield> [limit=<integer>]
```

- Multivalue command
- Takes a multivalue field and creates a separate event for each value; related field values are copied from the original event
 - <mvfield> is a multivalue field
 - Assigning an <integer> to **limit** determines the number of new events created from a multivalue field
- mvexpand is a distributable streaming command

field	mvfield
value1	mv-value1 mv-value2

```
...| mvexpand mvfield
```

field	mvfield
value1	mv-value1
value1	mv-value2

mvexpand Example

```
index=systems sourcetype=system_info
| rename ROOT_USERS{} as users, SYSTEM{} as system
| table system users
| mvexpand users
```

system	users
nix_system_idx_1	Duke
nix_system_idx_1	Camilian
nix_system_idx_1	Admin
nix_system_idx_1	Cardozo
nix_system_idx_4	Admin
nix_system_idx_4	Duke

Original event data that has been copied to each new event

Values of **users** have been split up into their own events

At this point, **users** is a multivalue field

system	users
nix_system_idx_1	Duke Camilian Admin Cardozo
nix_system_idx_4	Admin Duke Zed Camilian

Better Together: mvexpand + fields

- Remember: new events are created in memory, not in the index
- It is recommended to use **fields** before **mvexpand** to retain only the necessary fields for the search
 - Conserves memory usage and improves performance

```
index=systems sourcetype=system_info
| rename SYSTEM{} as system, CPU_CORES{}.core as core,
CPU_CORES{}.core_percent_used as percent_used
| table system core percent_used
| eval zip_percent_used=mvzip(core, percent_used, ": ")
| fields system zip_percent_used
| mvexpand zip_percent_used
```

This search has completed and has returned **1,037,232** results by scanning **86,436** events in **4.07 seconds**

```
index=systems sourcetype=system_info
| rename SYSTEM{} as system, CPU_CORES{}.core as core,
CPU_CORES{}.core_percent_used as percent_used
| table system core percent_used
| eval zip_percent_used=mvzip(core, percent_used, ": ")
| mvexpand zip_percent_used
```

This search has completed and has returned **1,037,160** results by scanning **86,430** events in **6.911 seconds**

Analyze Multivalue Data Lab Exercise

Time: 10 minutes

Tasks:

- Display information about AWS system data by completing a search with multivalue `eval` and `stats` functions
- Challenge: Modify the previous search so that `_number` is removed from each value of `CPU_CORES{}.core`

Wrap-up Slides

Community

- Splunk Community Portal
community.splunk.com
 - Answers
 - Discussions
 - Splunk Trust
 - User Groups
 - Ideas
- Splunk Blogs
splunk.com/blog/
- Splunk Apps
splunkbase.com
- Splunk Dev Google Group
groups.google.com/forum/#!forum/splunkdev
- Splunk Docs on Twitter
twitter.com/splunkdocs
- Splunk Dev on Twitter
twitter.com/splunkdev
- Splunk Live!
splunklive.splunk.com
- .conf
conf.splunk.com

Support Programs

- Web

- Documentation: dev.splunk.com and docs.splunk.com
- Wiki: wiki.splunk.com

- Splunk Lantern

Guidance from Splunk experts

- lantern.splunk.com

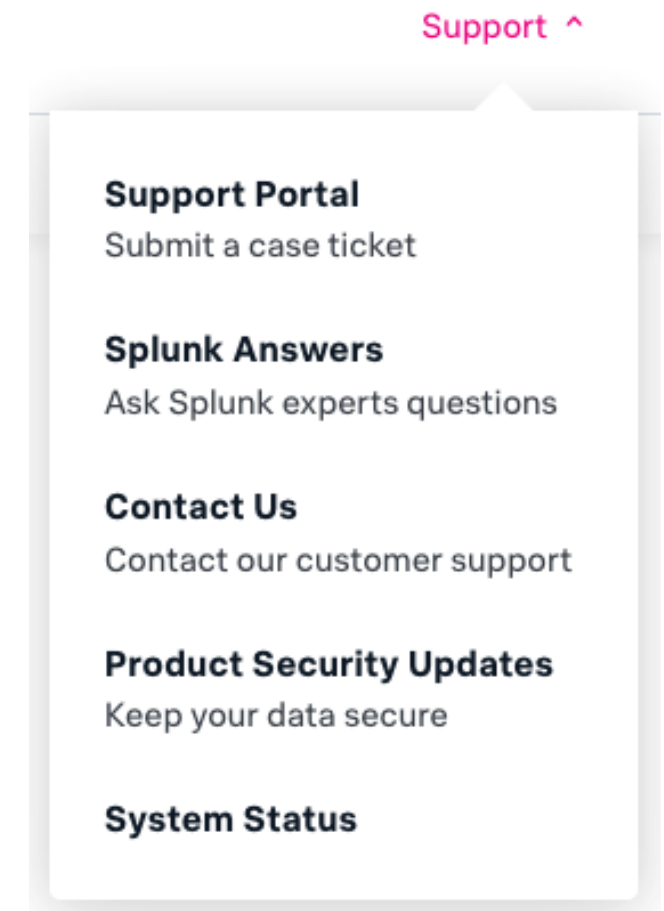
- Global Support

Support for critical issues, a dedicated resource to manage your account – 24 x 7 x 365

- Web: splunk.com/index.php/submit_issue

- Enterprise, Cloud, ITSI, Security Support

- Web: splunk.com/en_us/about-splunk/contact-us.html#tabs/customersupport
- Phone: (855) SPLUNK-S or (855) 775-8657



Learning Paths (cont.)

Knowledge Manager - Recommended Courses

Free eLearning courses are in [blue](#) and courses with an * are present in both learning paths.

- [What is Splunk *](#)
- [Introduction to Splunk *](#)
- [Using Fields *](#)
- [Introduction to Knowledge Objects](#)
- [Creating Knowledge Objects](#)
- [Creating Field Extractions](#)
- [Enriching Data with Lookups](#)
- [Data Models](#)
- [Introduction to Dashboards](#)
- [Dynamic Dashboards](#)
- [Using Choropleth](#)
- [Search Optimization *](#)

Learning Paths

Search Expert - Recommended Courses

Free eLearning courses are in [blue](#) and courses with an * are present in both learning paths.


- [What is Splunk *](#)
- [Introduction to Splunk *](#)
- [Using Fields *](#)
- [Scheduling Reports and Alerts](#)
- [Visualizations](#)
- Statistical Processing
- Working with Time
- Comparing Values
- Result Modification
- Leveraging Lookups and Subsearches
- Correlation Analysis
- [Search Under the Hood](#)
- Multivalue Fields
- Search Optimization *

Splunk Certification

Offerings & Requirements

Splunk Core and Beyond

Regardless of which Splunk product you use, it all starts with Splunk Core



Splunk Core Certified Advanced Power User

This certification demonstrates an individual's ability to generate complex searches, reports, and dashboards with Splunk's core software to get the most out of their data



Prerequisite Certification(s):

- [Splunk Core Certified Power User](#)

Prerequisite Course(s):

- None



Splunk Core Certified Advanced Power User Exam


Time to [study](#)! We suggest candidates looking to prepare for this exam complete Fundamentals 3, Creating Dashboards, and Advanced Searching & Reporting **or** the following courses:

- Using Fields
- Working with Time
- Comparing Values
- Result Modification
- Leveraging Lookups and Subsearches
- Correlation Analysis
- Search Under the Hood
- Multivalue Fields
- Search Optimization
- Creating Field Extractions
- Enriching Data with Lookups
- Data Models
- Using Choropleth
- Introduction to Dashboards
- Dynamic Dashboards

See [here](#) for registration assistance.



Congratulations! You are a...



Recommended Next Steps

- [Splunk Enterprise Certified Admin](#)
- [Splunk Cloud Certified Admin](#)

Thank You

