Enriching Data with Lookups

Document Usage Guidelines

- Should be used only for enrolled students
- Not meant to be a self-paced document, an instructor is needed
- Lab Exercise slides reference the hands-on lab exercise guide
- Do not distribute

Course Goals

- Define lookups
- Identify types of lookups
- Create a lookup
- Define a geospatial lookup
- Use an external lookup
- Define a KV Store lookup

Course Outline

- What is a Lookup?
- Create a Lookup
- Geospatial Lookups
- External Lookups
- KV Store Lookups
- Best Practices for Lookups

What is a Lookup?

Topic Objectives

- Define a lookup and the default lookup types
- Where lookups fall in the search-time operation sequence

What is a Lookup?

- **Lookups** provide enrichment to your event data by appending fields from another data source (i.e. lookup output fields)
- Splunk provides four types of lookups by default

Lookup Type	Description
File-based	Populates your events with fields pulled from CSV files
External	Uses Python scripts or binary executables to append data
KV Store	Accesses key value pairs from a KV Store collection
Geospatial	References a KMZ or KML file

Search-time Operation Sequence

Search-time operations are always applied in the same order when generating the knowledge objects

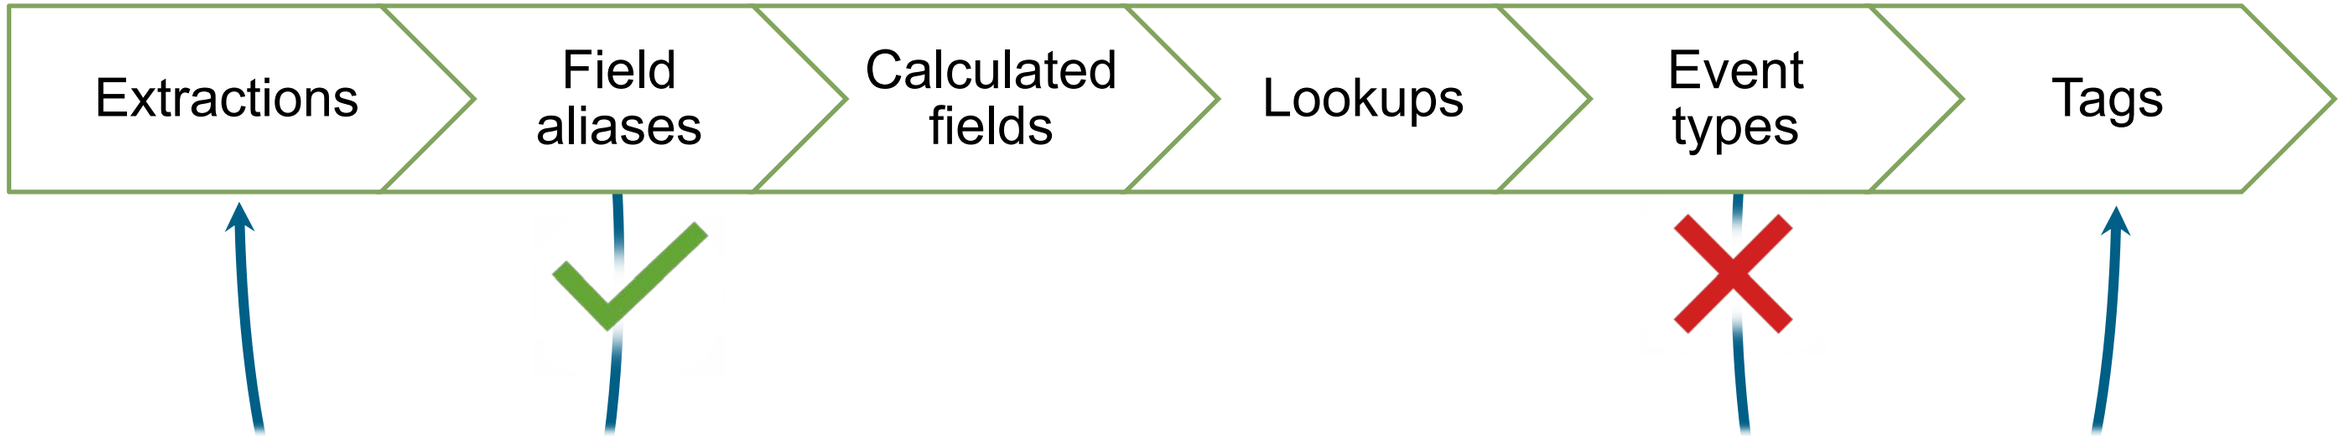

Each operation can reference knowledge objects derived from operations that **precede** them

No operation can reference knowledge objects derived by operations that **follow** them

Create a Lookup

Topic Objectives

- Describe lookups at search time
- Use file-based lookups
- Examine a CSV lookup file
- Create a lookup
  - Upload a lookup table file
  - Define a lookup
  - Configure time-based lookup
  - Apply advanced lookup options
- Create and use an automatic lookup at search

Lookups at Search Time

- Sometimes static (or relatively unchanging) data is required for searches, but isn't available in the raw event data
- Lookups pull such data from standalone files at search time and add it to search results as field values

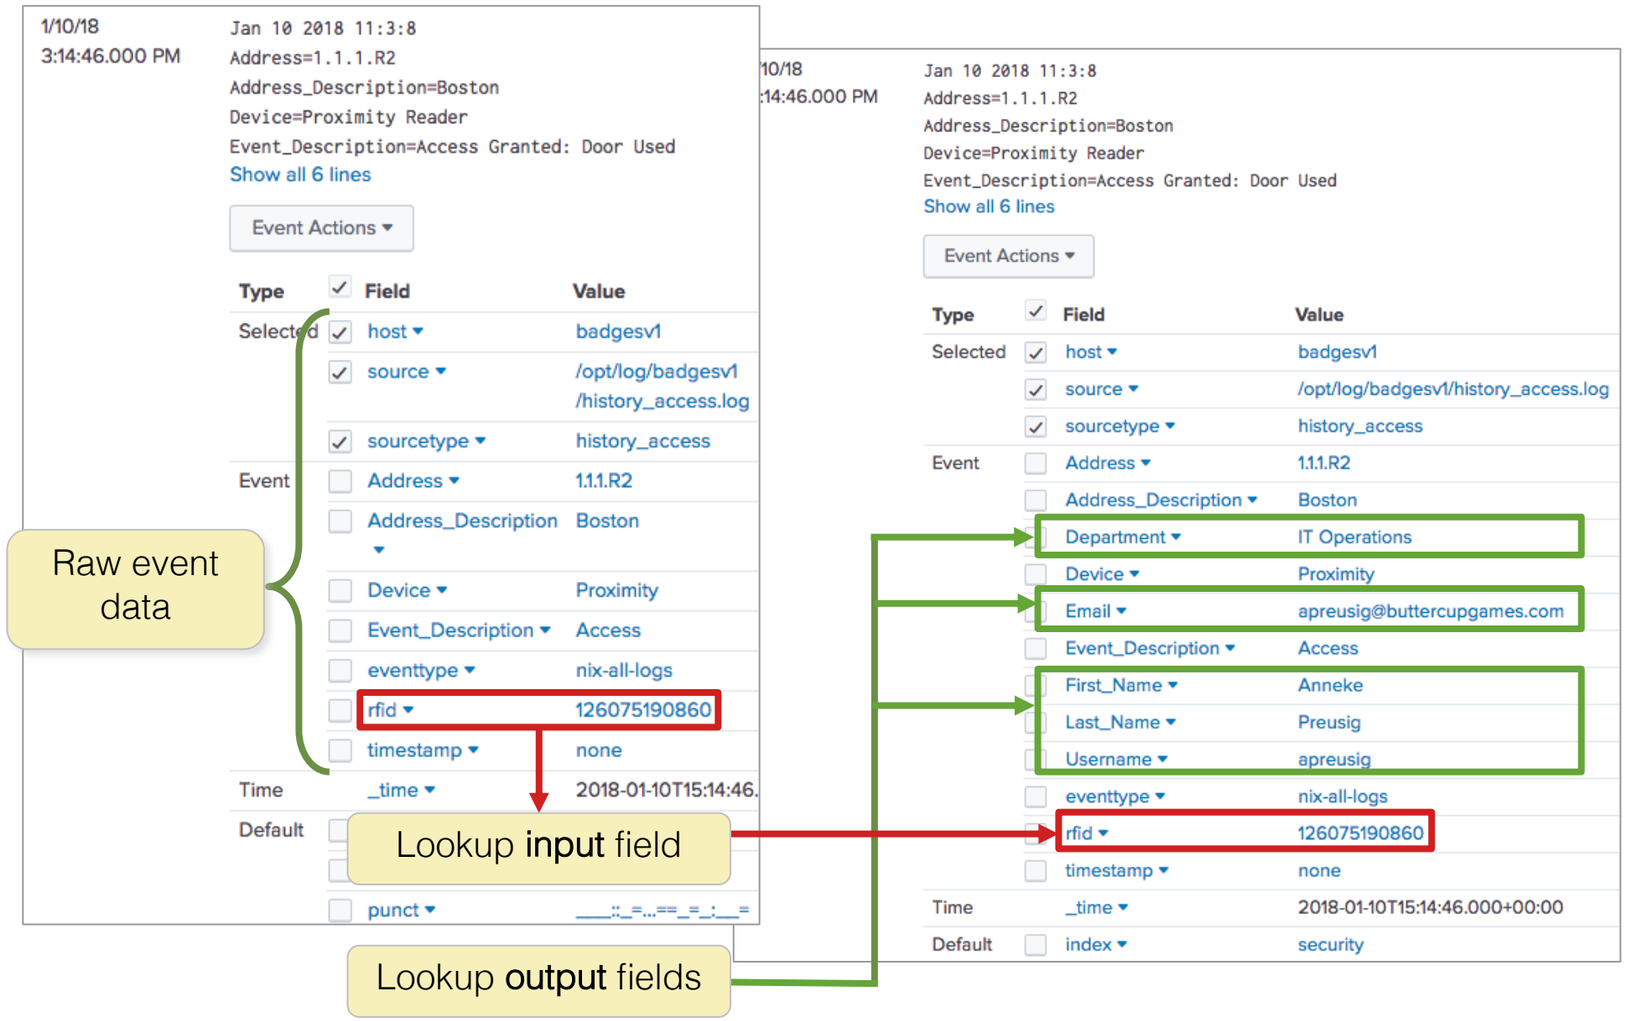

Use a File-Based Lookup

- Lookups allow you to add more fields to your events such as:
  - Descriptions for HTTP status codes (“Not Found”, “Service Unavailable”)
  - Sale prices for products
  - Usernames, IP addresses, and workstation IDs associated with RFIDs
- After a lookup is invoked, lookup fields appear in the Fields sidebar and can be used in searches and reports
- Lookups can be invoked by the **`lookup`** command or configured to run automatically
- Lookup field values are case sensitive by default

A Sample CSV Lookup File

First row represents field names (header)

Input field: The `productId` field exists in both `access_combined` events and `.csv` file

Output fields: All other fields in the `.csv` file are searchable after the lookup is defined

```
 GNU nano 2.3.1                File: products.csv

productId,product_name,categoryId,price,sale_price,Code
DB-SG-G01,Mediocre Kingdoms,STRATEGY,24.99,19.99,A
DC-SG-G02,Dream Crusher,STRATEGY,39.99,24.99,B
FS-SG-G03,Final Sequel,STRATEGY,24.99,16.99,C
WC-SH-G04,World of Cheese,SHOOTER,24.99,19.99,D
WC-SH-T02,World of Cheese Tee,TEE,9.99,6.99,E
PZ-SG-G05,Puppies vs. Zombies,STRATEGY,4.99,1.99,F
CU-PG-G06,Curling 2014,SPORTS,19.99,16.99,G
MB-AG-G07,Manganiello Bros.,ARCADE,39.99,24.99,H
MB-AG-T01,Manganiello Bros. Tee,TEE,9.99,6.99,I
FI-AG-G08,Orvil the Wolverine,ARCADE,39.99,24.99,J
BS-AG-G09,Benign Space Debris,ARCADE,24.99,19.99,K
SC-MG-G10,SIM Cubicle,SIMULATION,19.99,16.99,L
WC-SH-A01,Holy Blade of Gouda,ACCESSORIES,5.99,2.99,M
WC-SH-A02,Fire Resistance Suit of Provolone,ACCESSORIES,3.99,1.99,N
```

Create a Lookup

1. Upload the file required for the lookup
2. Define the lookup type
3. Optionally, configure the lookup to run automatically

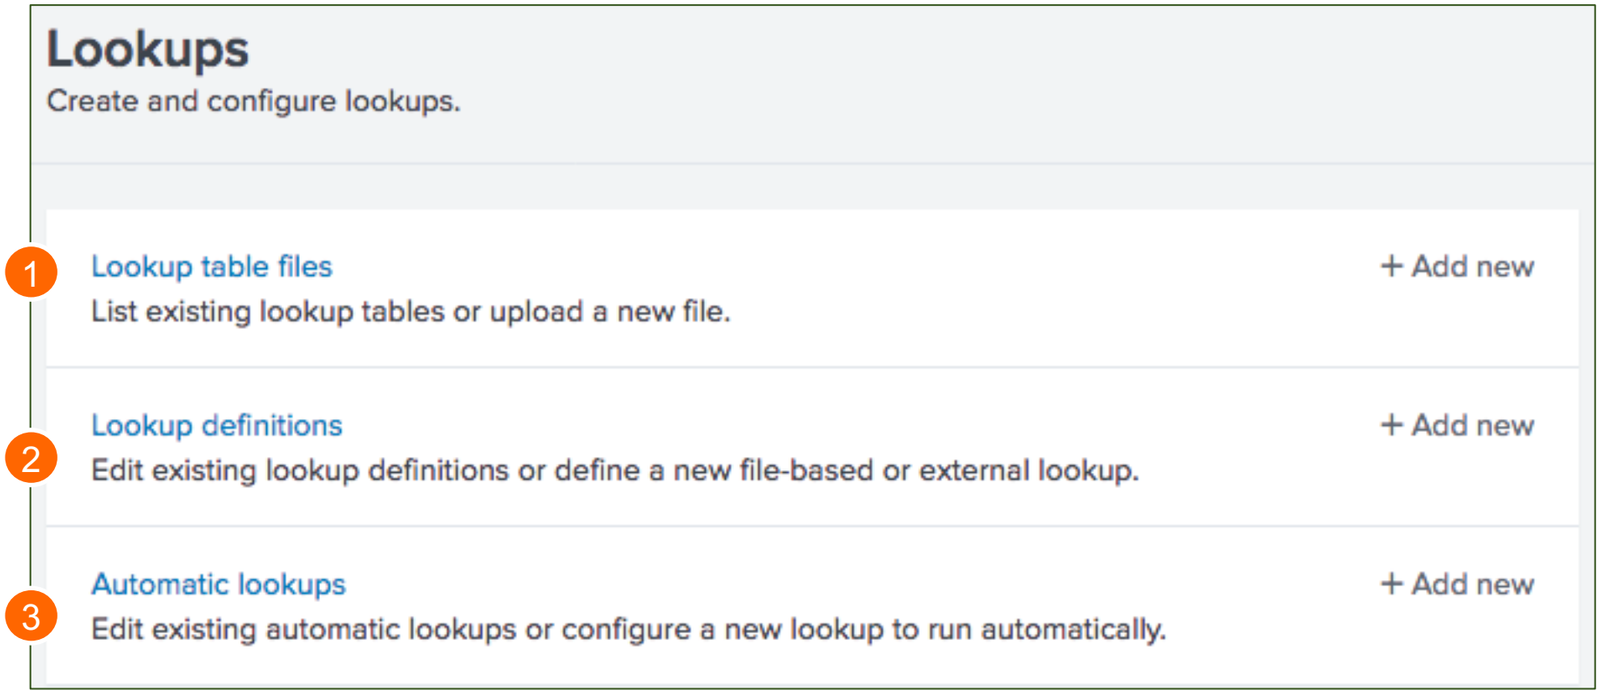

Add a New Lookup Table File

Settings > Lookups > Lookup table files > 1 **New Lookup Table File**

1. Click **New Lookup Table File**
2. Select a **Destination app**
3. Browse and select the file to use for the lookup table
4. Enter a name for the lookup file
5. Save

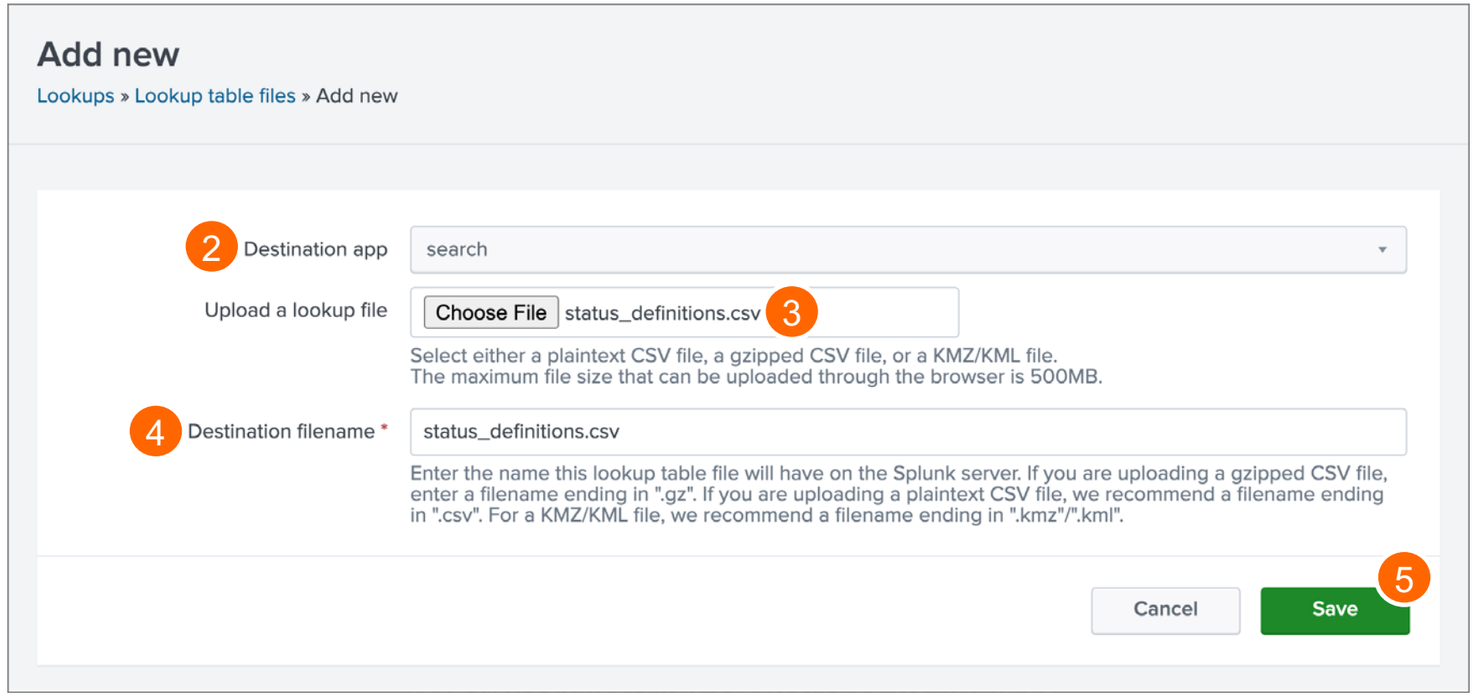

Note

For file-based lookups, the lookup should be a CSV file or a gzipped CSV file.

Create a Lookup Definition

Settings > Lookups > Lookup definitions> 1 **New Lookup Definition**

1. Click **New Lookup Definition**
2. Select a **Destination app**
3. **Name** the Lookup definition
4. For **Type**, select **File-based**
5. Browse and select the file to use for the lookup table
6. Save

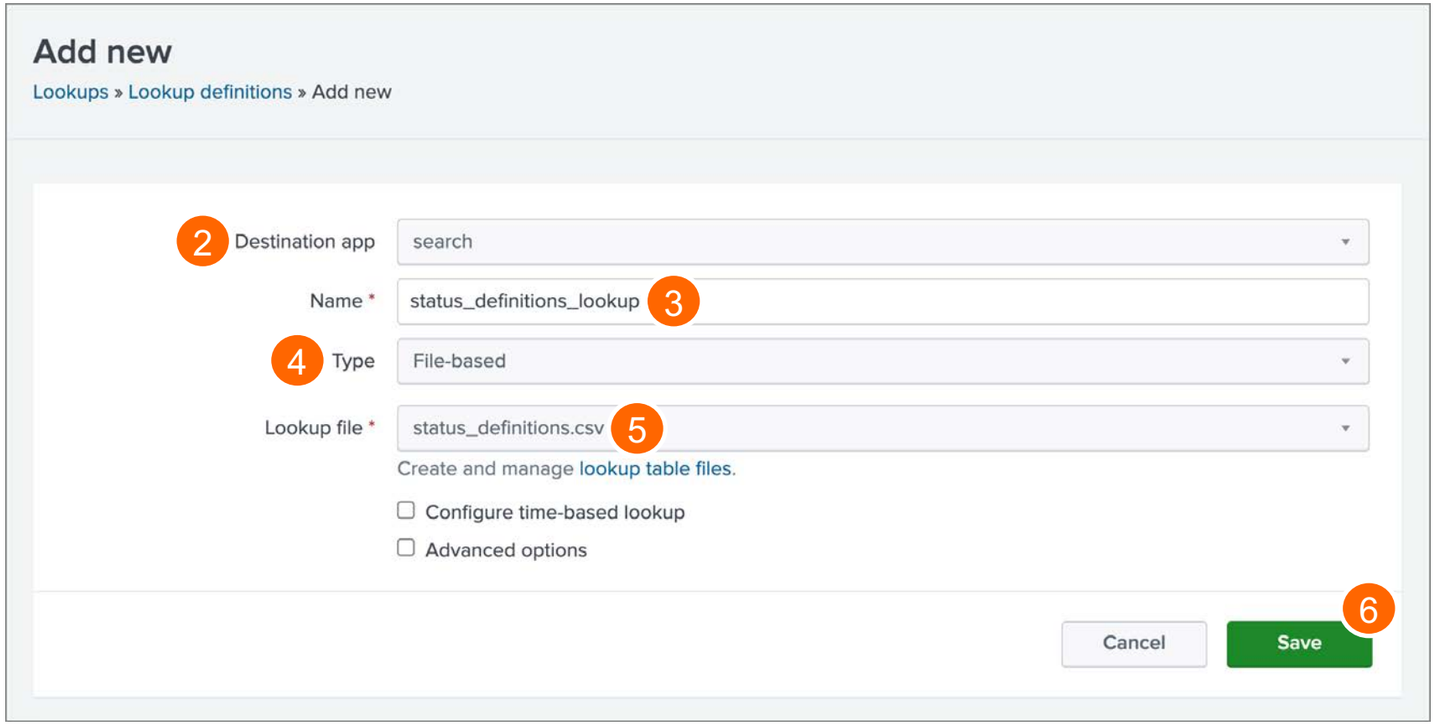

Lookup is now listed in **Settings** > **Lookups** > **Lookup definitions**

Name	Type	Supported fields	Lookup file
dnslookup	external	clienthost,clientip	
employee_lookup	file	RFID,FIRSTNAME,LASTNAME,USERNAME,EMAIL,DEPT,LOCATION,IP,HOSTNAME,SPLUNKROLE	employees.csv

Configure Time-based Lookup

1. Field in the record that represents the timestamp
2. Specifies the strptime() format of the time_field attribute (%s.%Q represents Unix epoch time in seconds and milliseconds)
3. Minimum amount of time in seconds that an event timestamp can be later than the record timestamp
4. Maximum amount of time in seconds that an event timestamp can be later than the record timestamp

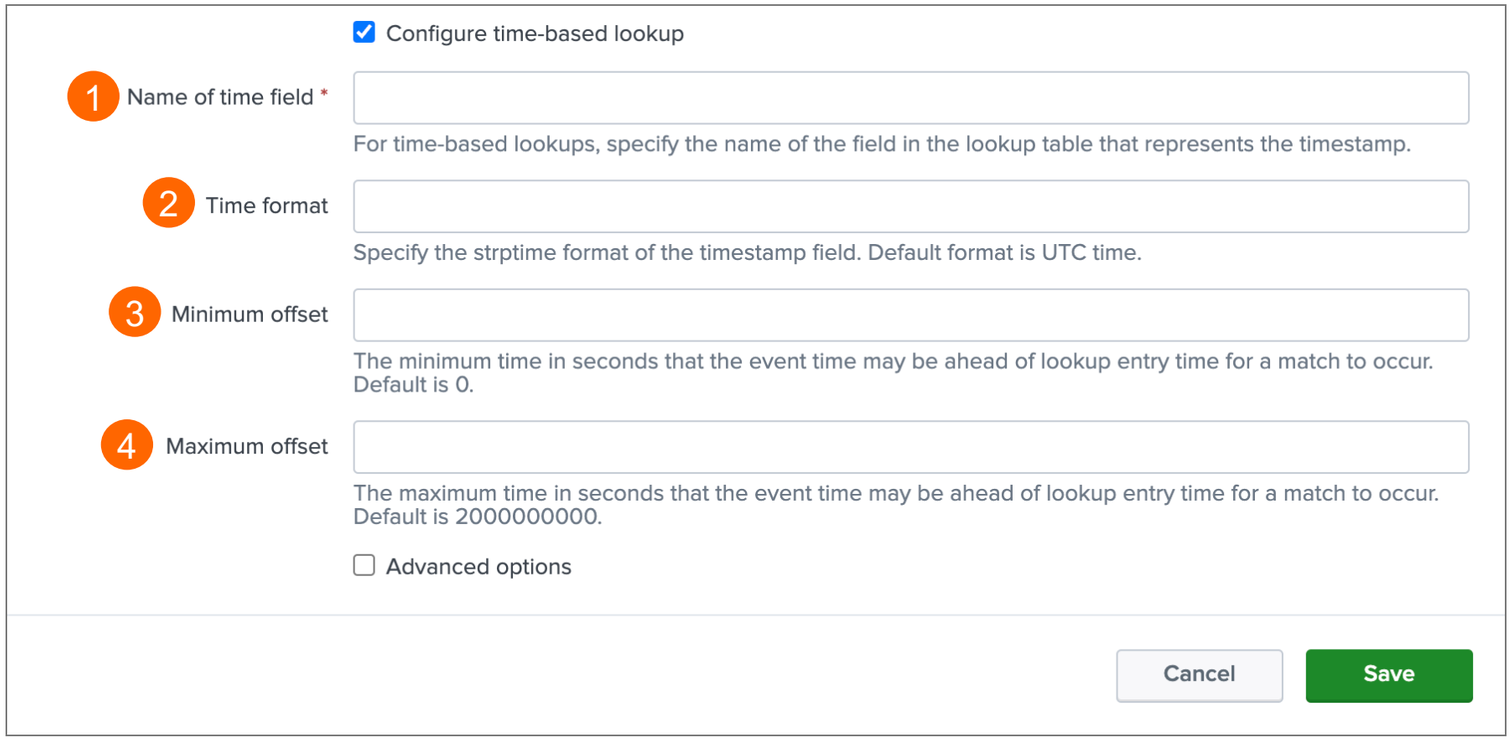

Apply Advanced Lookup Options

Minimum matches: min # of matches for each input lookup value

Default matches: value to output when fewer than the minimum matches are returned for a given input; defaults to an empty string

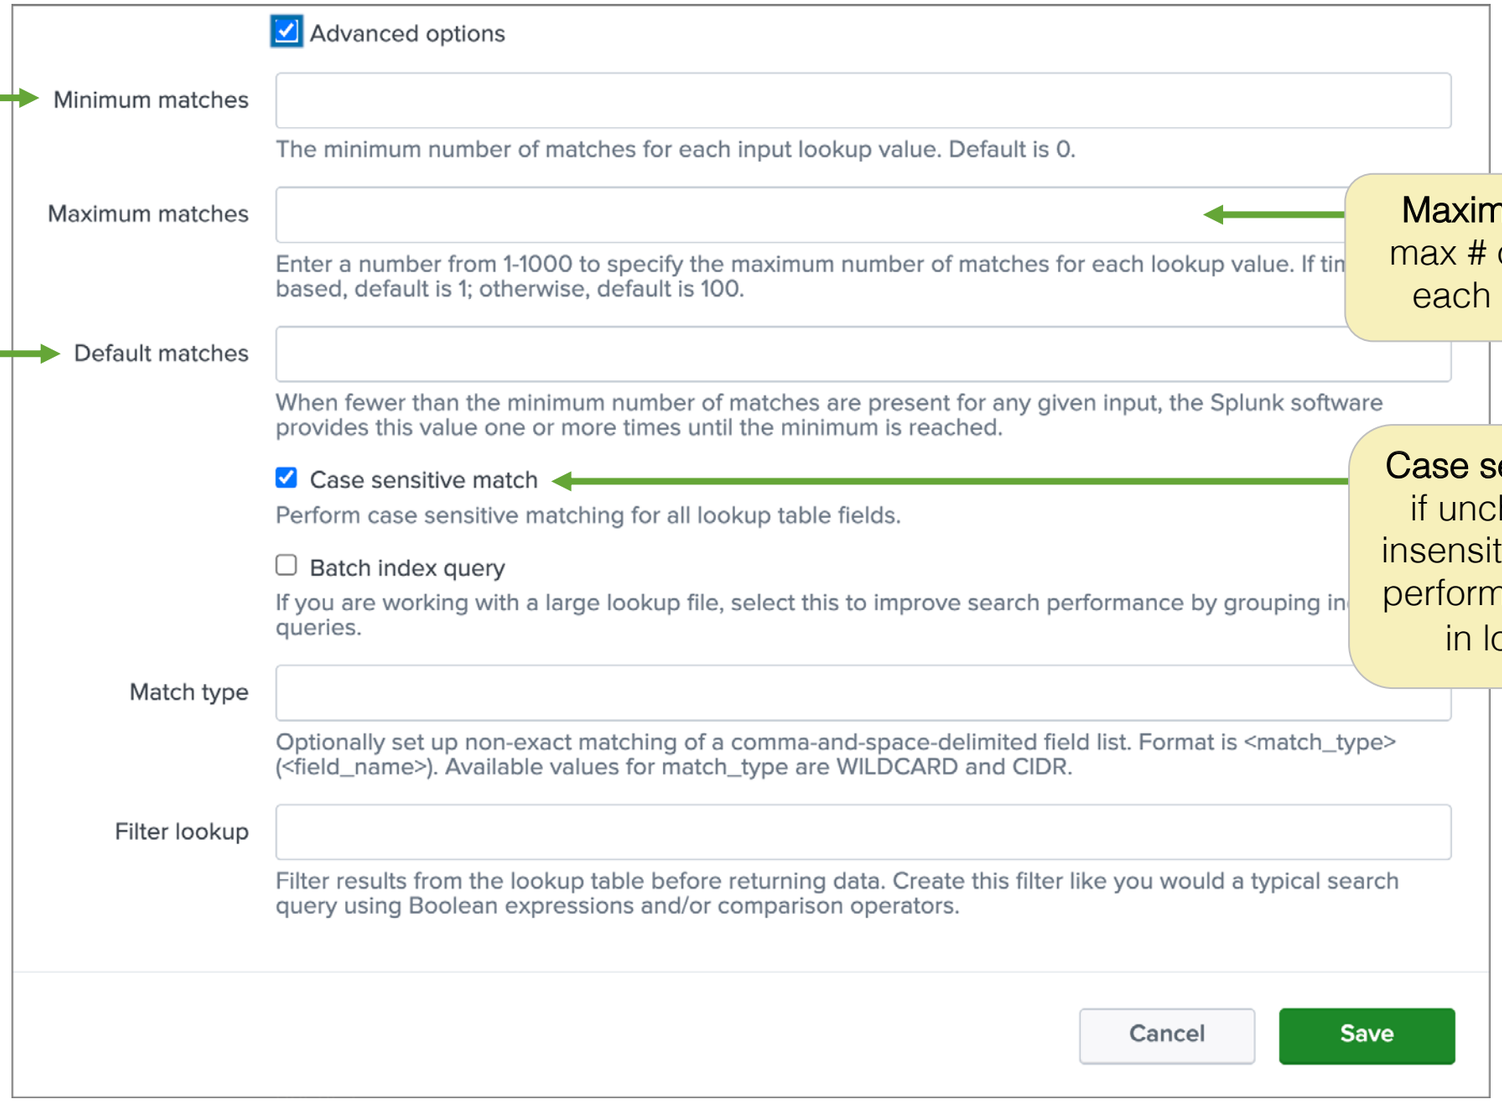

Maximum matches: max # of matches for each lookup value

Case sensitive match: if unchecked, case insensitive matching is performed for all fields in lookup table

 | 12 August 2022

Apply Advanced Lookup Options (cont.)

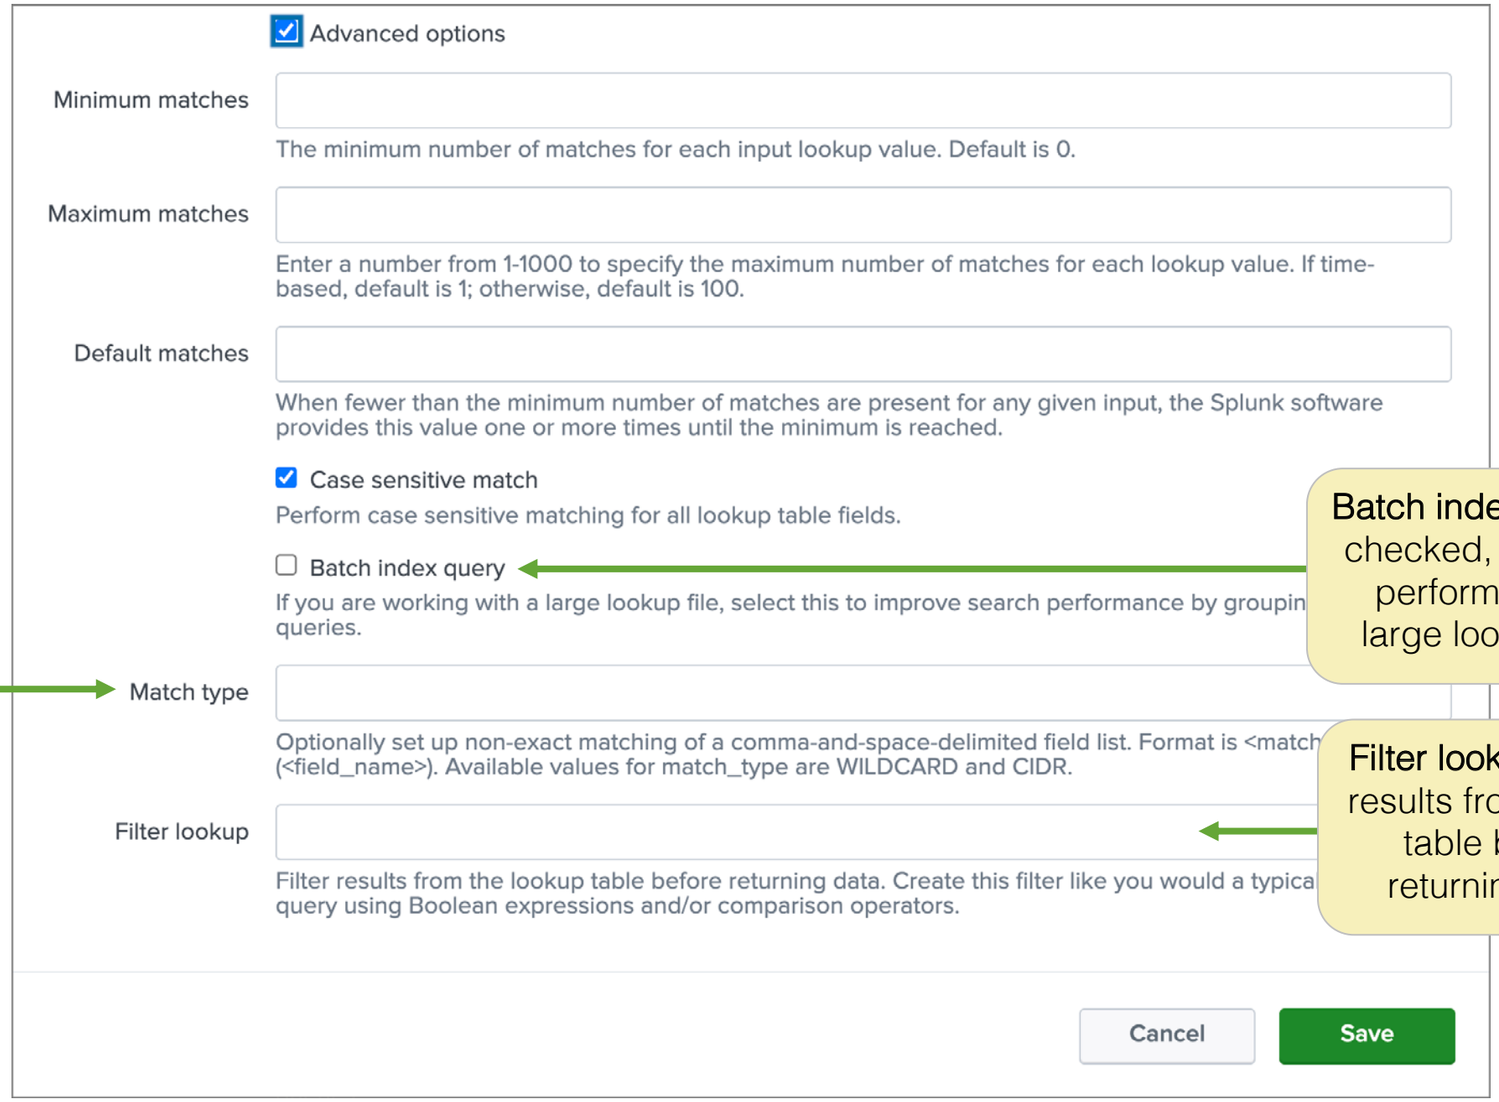

Match type: allows for non-exact matching using the match types WILDCARD and CIDR

Batch index query: if checked, improves performance of large lookup files

Filter lookup: filters results from lookup table before returning data

 | 12 August 2022

Create an Automatic Lookup

Settings > Lookups > Automatic lookups > 1 **New Automatic Lookup**

1. Click **New Automatic Lookup**
2. Select a **Destination app**
3. **Name** the automatic lookup
4. Select the **Lookup table** definition
5. Select host, source, or sourcetype to apply to the lookup and specify the name

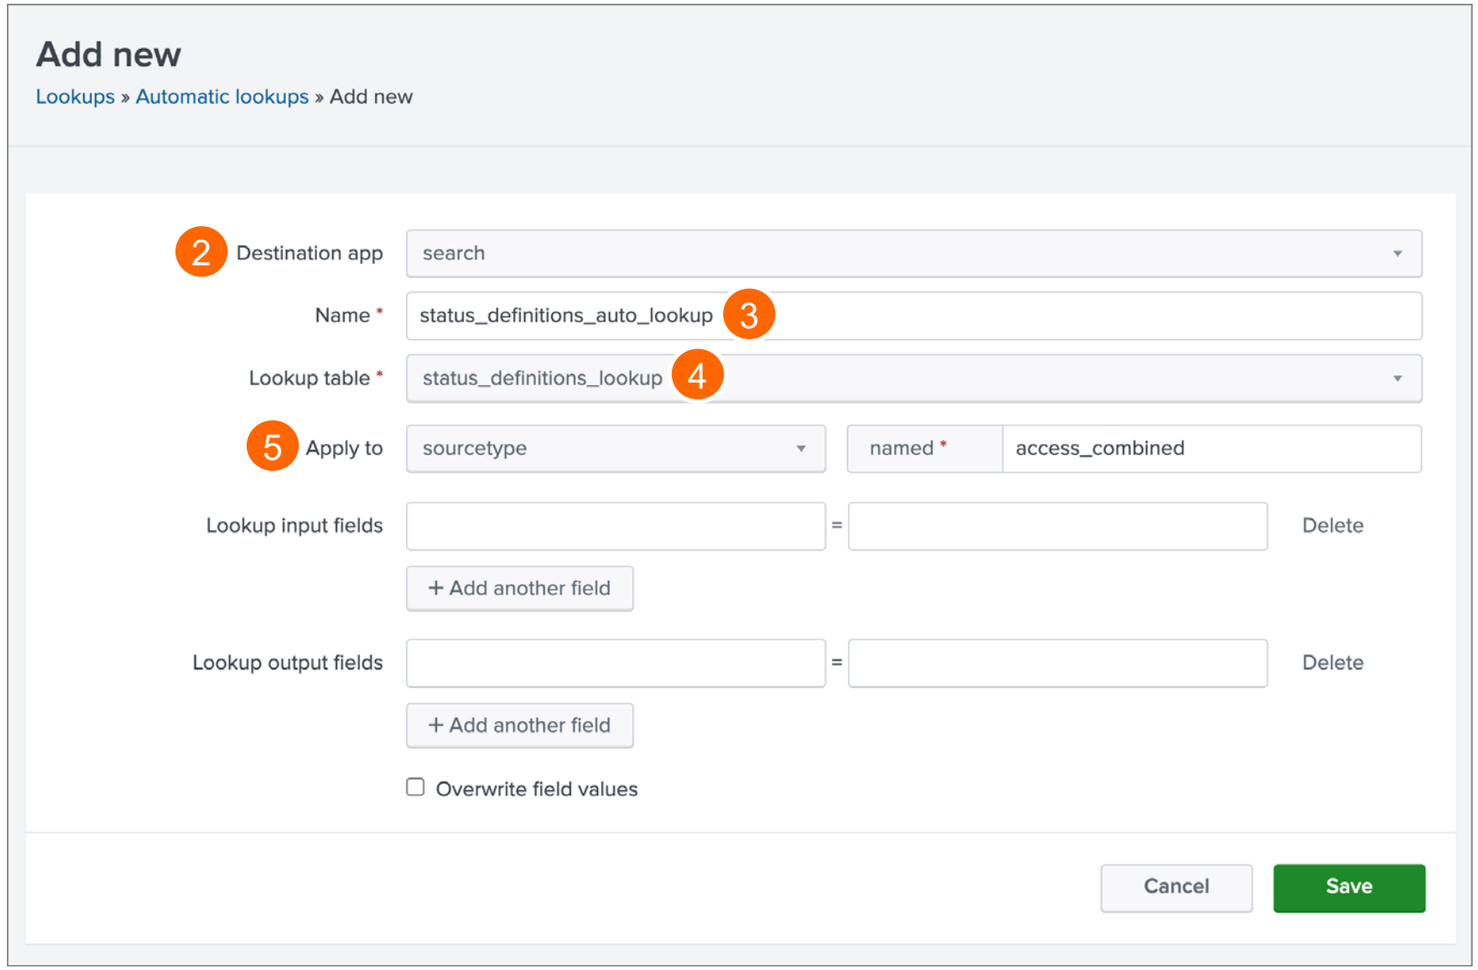

Create an Automatic Lookup (cont.)

Settings > Lookups > Automatic lookups > **New Automatic Lookup**

6. Define the **Lookup input fields**: the field(s) that exist in your events that you are relating to the lookup table
   - A Column name in CSV
   - B Field name in events
7. Define **Lookup output fields**
   - C Field name in lookup table
   - D Name you want displayed in results, otherwise column name from CSV is inherited
8. Save

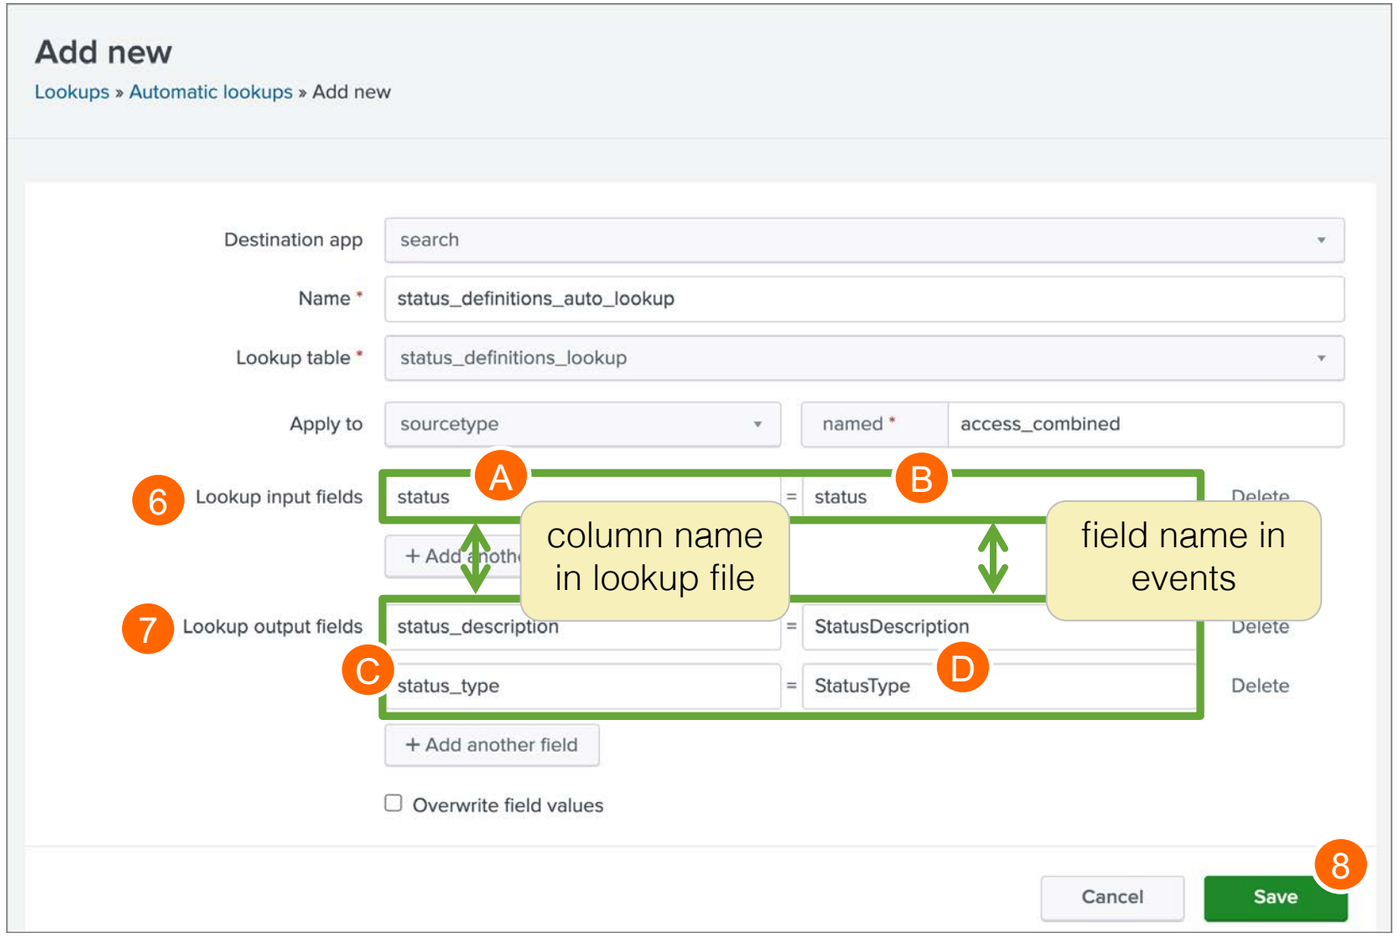

Create an Automatic Lookup (cont.)

- Lookup is now listed in **Lookups** > **Automatic lookups**
- Automatic lookups are applied to all searches at search time

Name	Lookup	Owner	App	Sharing	Status	Actions
access_combined : LOOKUP-status_definitions_auto_lookup	status_definitions_lookup status AS status OUTPUTNEW status_description AS StatusDescription status_type AS StatusType	poweruser	search	Private | Permissions	Enabled	Clone | Move | Delete

Use an Automatic Lookup in Search

To use an automatic lookup, specify the output fields in your search

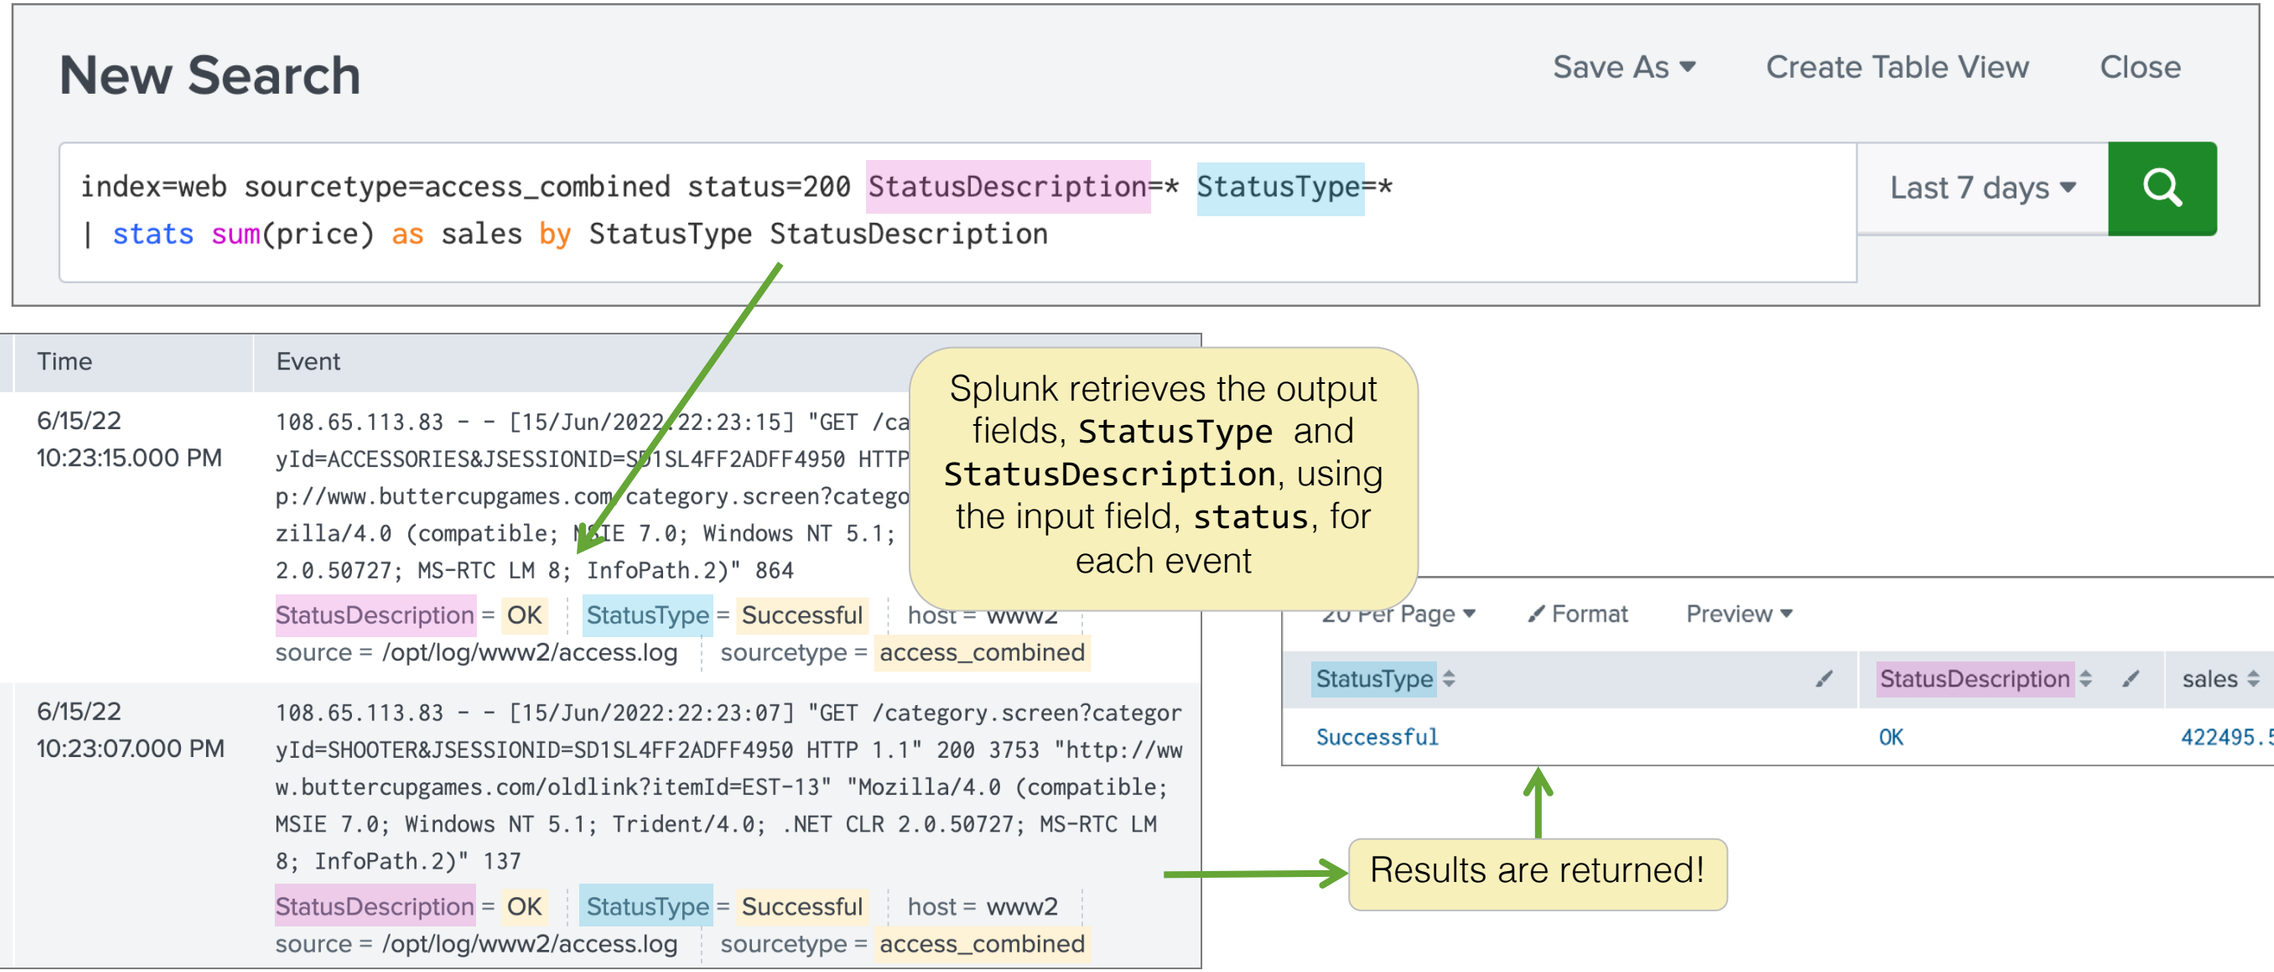

Create a Lookup Lab Exercise

Time: 20-25 minutes

Tasks:

- Add lookup table files to your search environment
- Create a lookup definitions
- Create an automatic lookup
- Verify your automatic lookup is working in search

Geospatial Lookups

Topic Objectives

- Describe the use of geospatial lookups
- Examine `KML/KMZ` geospatial lookup files
- Add a geospatial lookup file
- Define a geospatial lookup

Geospatial Lookups

- Matches region names in your events to region names in lookup and outputs fields with corresponding geographic feature info
- Location coordinate ranges are provided by geographic feature collections: `.KML` and `.KMZ` files
- Geospatial lookups can be invoked in searches to generate choropleth map visualizations
- Splunk ships with two geospatial lookup files:
  - `geo_us_states`
  - `geo_countries`

Geospatial Lookup Files

- Provides geographic feature information used to define a geospatial lookup
  - **KML**: a type of XML file
  - **KMZ**: a zipped KML file
- Rely on polygons which are closed shapes that start and end at the same coordinate
- Many are available online or can be created from scratch using software such as Google Earth

geo_countries.kml

```
<?xml version="1.0" encoding="utf-8" ?>
<kml xmlns="http://www.opengis.net/kml/2.2">
<Document id="root_doc">
<Schema name="countries" id="countries">
         <SimpleField name="Name" type="string"></SimpleField>
         <SimpleField name="ISO2" type="string"></SimpleField>
         <SimpleField name="ISO3" type="string"></SimpleField>
         <SimpleField name="REGION_WB" type="string"></SimpleField>
         <SimpleField name="REGION_UN" type="string"></SimpleField>
         <SimpleField name="SUBREGION" type="string"></SimpleField>
         <SimpleField name="CONTINENT" type="string"></SimpleField>
</Schema>
<Folder><name>countries</name>
  <Placemark>
         <name>Aruba</name>
         <Style><LineStyle><color>ff0000ff</color></LineStyle><PolyStyle><fill>0</fill></
PolyStyle></Style>
         <ExtendedData><SchemaData schemaUrl="#countries">
                  <SimpleData name="ISO2">AW</SimpleData>
                  <SimpleData name="ISO3">ABW</SimpleData>
                  <SimpleData name="REGION_WB">Latin America &amp; Caribbean</SimpleData>
                  <SimpleData name="REGION_UN">Americas</SimpleData>
                  <SimpleData name="SUBREGION">Caribbean</SimpleData>
                  <SimpleData name="CONTINENT">North America</SimpleData>
         </SchemaData></ExtendedData>
      <Polygon><outerBoundaryIs><LinearRing><coordinates>-
69.996937628999916,12.577582098000036 -69.924672003999945,12.519232489000046 -
69.880197719999842,12.453558661000045 -69.888091600999928,12.417669989000046 -
69.930531378999888,12.425970770000035 -69.945139126999919,12.44037506700009 -
69.924672003999945,12.447211005000014 -70.058094855999883,12.537176825000088 -
70.048736131999931,12.583726304000024 -70.061105923999975,12.625392971000068 -
70.048736131999931,12.632147528000104 -
69.996937628999916,12.577582098000036</coordinates></LinearRing></outerBoundaryIs></
Polygon>
  </Placemark>
```

`geo_countries` content for the island country of Aruba. The `Polygon` tag (highlighted) contains the coordinates Splunk uses to define its choropleth map data.

Add a Geospatial Lookup Table File

Settings > Lookups > Lookup table files > **New Lookup Table File**

1. Select a **Destination app**
2. Browse and select the `.kmz` or `.kml` file to use for the lookup table
3. Enter a name for the lookup file
4. Save

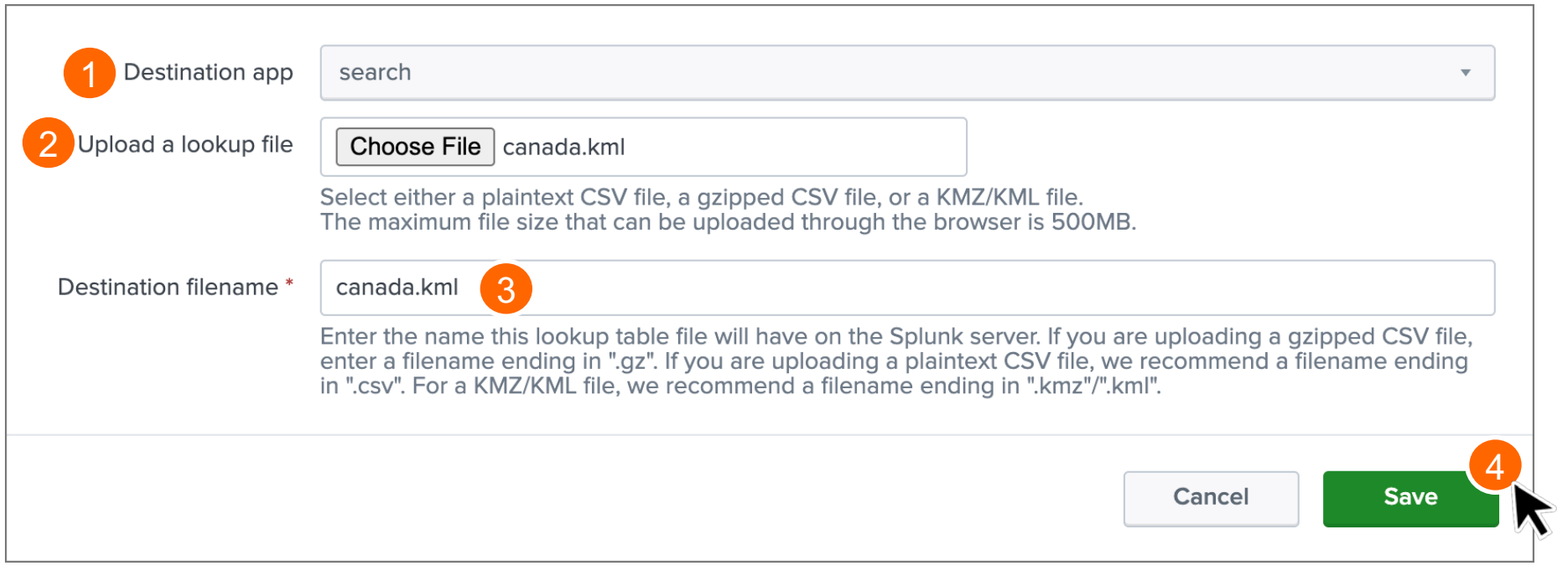

Define a Geospatial Lookup

Settings > Lookups > Lookup definitions > **New Lookup Definition**

1. Select a **Destination app**
2. **Name** the lookup definition
3. Change **Type** to **Geospatial**
4. Select the **Lookup file** from the drop-down list
5. Save

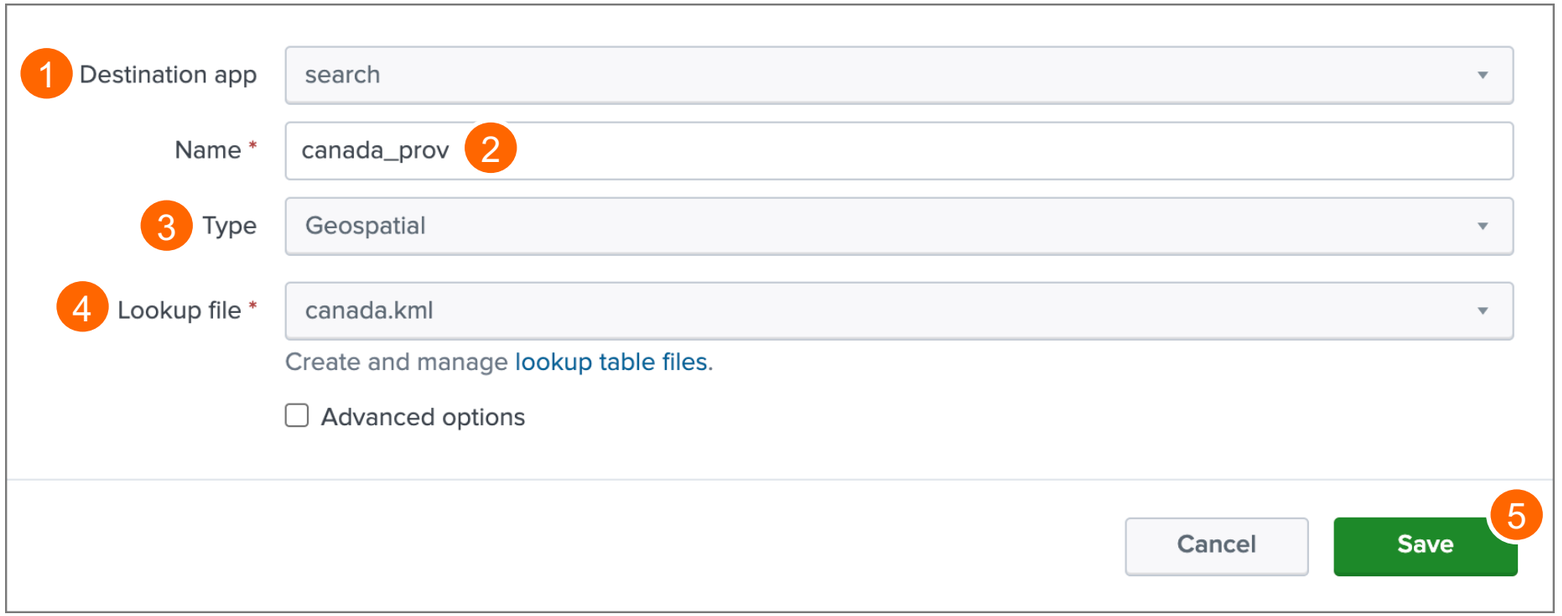

External Lookups

Topic Objectives

- Define the use of external lookups
- Examine an **`external_lookup.py`** lookup script
- Configure external lookups

What are External Lookups

- External lookups use scripts or executables to populate events with field values from an external source
- Often referred to as **scripted lookups**

```
index=web sourcetype=access_combined
| lookup dnslookup clientip
| stats count by clienthost
```

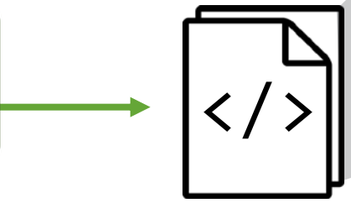

external_lookup.py

```
# Given an ip, return the host
def rlookup(ip):
    try:
        hostname, aliaslist, ipaddrlist = socket.gethostbyaddr(ip)
        return hostname
    except:
        return ''
```

Underlying script or executable

External source

clienthost	count
201-28-109-162.customer.tdatabrasil.net.br	5
208.red-88-12-32.staticip.rima-tde.net	8
72.246.204.221.adsl-pool.sx.cn	5
87-194-216-51.bethere.co.uk	17
88.218.220.60.adsl-pool.sx.cn	13
dynamicip-176-212-0-44.pppoe.kirov.ertelecom.ru	5
h-sys.biz	4

Manage an External Lookup Script

- Must be a Python script or binary executable
- Must be added to your Splunk deployment in either:
  - `$SPLUNK_HOME/etc/searchscripts`
  - `$SPLUNK_HOME/etc/apps/<app_name>/bin`

external_lookup.py

- Splunk ships with a sample script **external_lookup.py** in **$SPLUNK_HOME/etc/system/bin**

To use the sample script:

1. Move **external_lookup.py** script to appropriate directory
2. Create **dnslookup** definition as shown in next slide
3. Invoke the lookup using either:

```
...| lookup dnslookup clienthost
```

```
...| lookup dnslookup clientip
```

- Splunk passes values for **clienthost** into script and script returns **clientip** (or vice versa)
- Returned values are used to populate **clientip** or **clienthost** in the results

Note

The first step has already been completed in the lab environment.

Configure an External Lookup

Settings > Lookups > Lookup definitions > **New Lookup Definition**

1. Select **Destination app**
2. **Name** the lookup definition
3. Change **Type** to **External**
4. Enter script name and arguments passed to script
5. List all fields supported by the lookup
6. Save

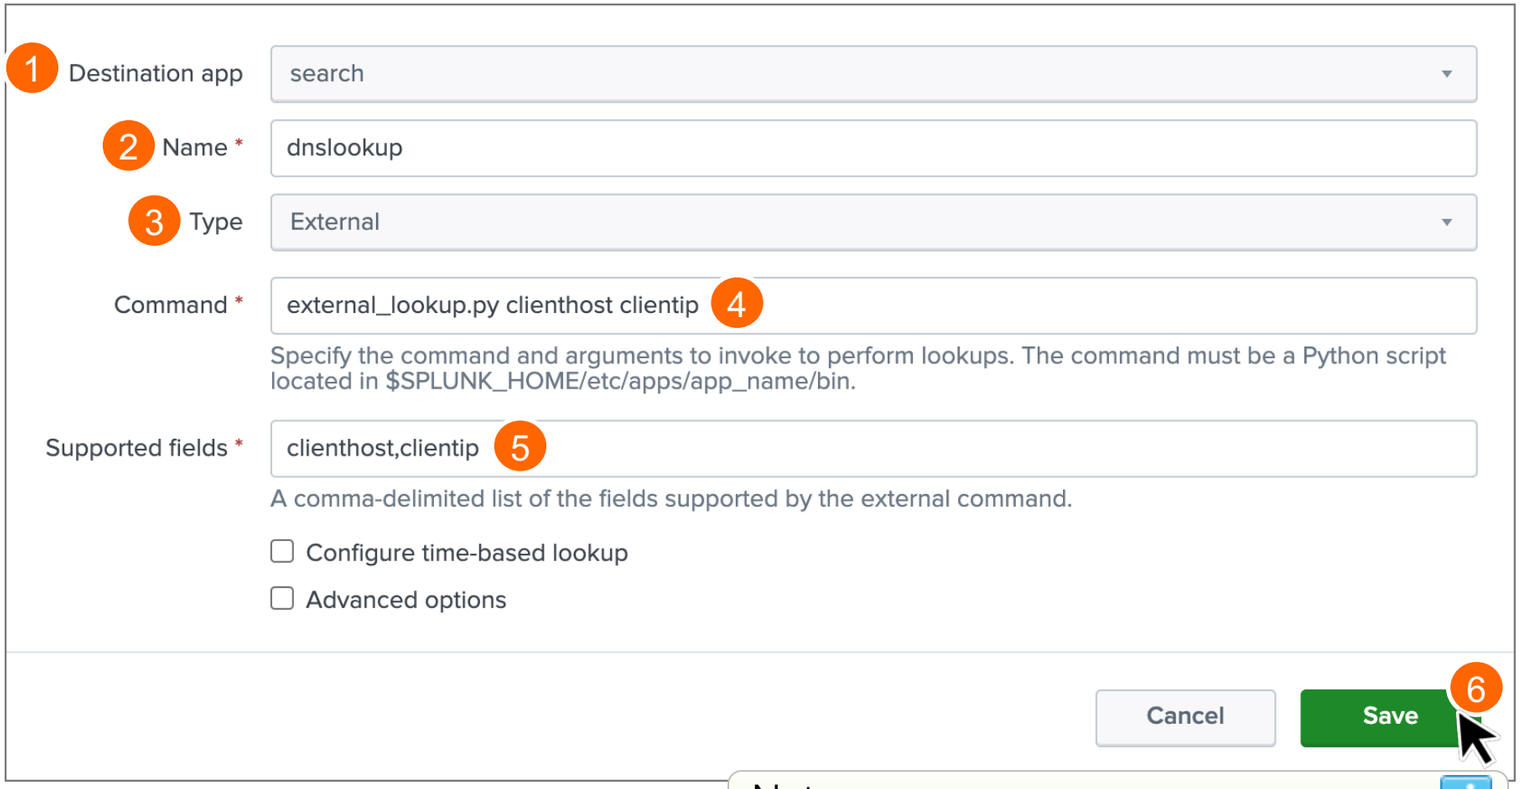

Note

The arguments passed to the script are the field headers from the input/output CSV files.

Geospatial & External Lookups Lab Exercise

Time: 10 minutes

Tasks:

- Upload and define a geospatial lookup and verify its contents in search
- Define an external lookup and use it in search

KV Store Lookups

Topic Objectives

- Define the use of KV Store lookups
- Identify the steps to set up a KV Store lookup
- Examine the KV Store lookups `collections.conf` file
- Create a KV Store lookup definition
- Identify the options for populating a KV store lookup
- Compare file-based CSV lookups to KV Store lookups

Use KV Store Lookups

- Instead of matching against values in a CSV file, you can also match against values in a **KV Store** (key value store)
- Use for large lookup tables or ones that are updated often
- KV Store saves and retrieves data in **collections** of key-value pairs
  - Similar to database tables in which each record has a unique key
  - Provides multiuser access locking so that multiple users can not edit the same record at the same time

Steps to Set up a KV Store Lookup

1. Add configuration stanzas to **`collections.conf`** (admin only)
2. Create KV Store definition
3. Populate the KV Store lookup with data using:
   - **`outputlookup`** command (admin and power user ability)
   - REST API (admin ability)
   - A front-end form (not discussed in this course)

Note

Once defined, the admin can share the KV Store collection with other apps and users

Examine the KV Store: `collections.conf`

An admin must add a stanza for each KV Store in the `collections.conf` file before a definition can be created

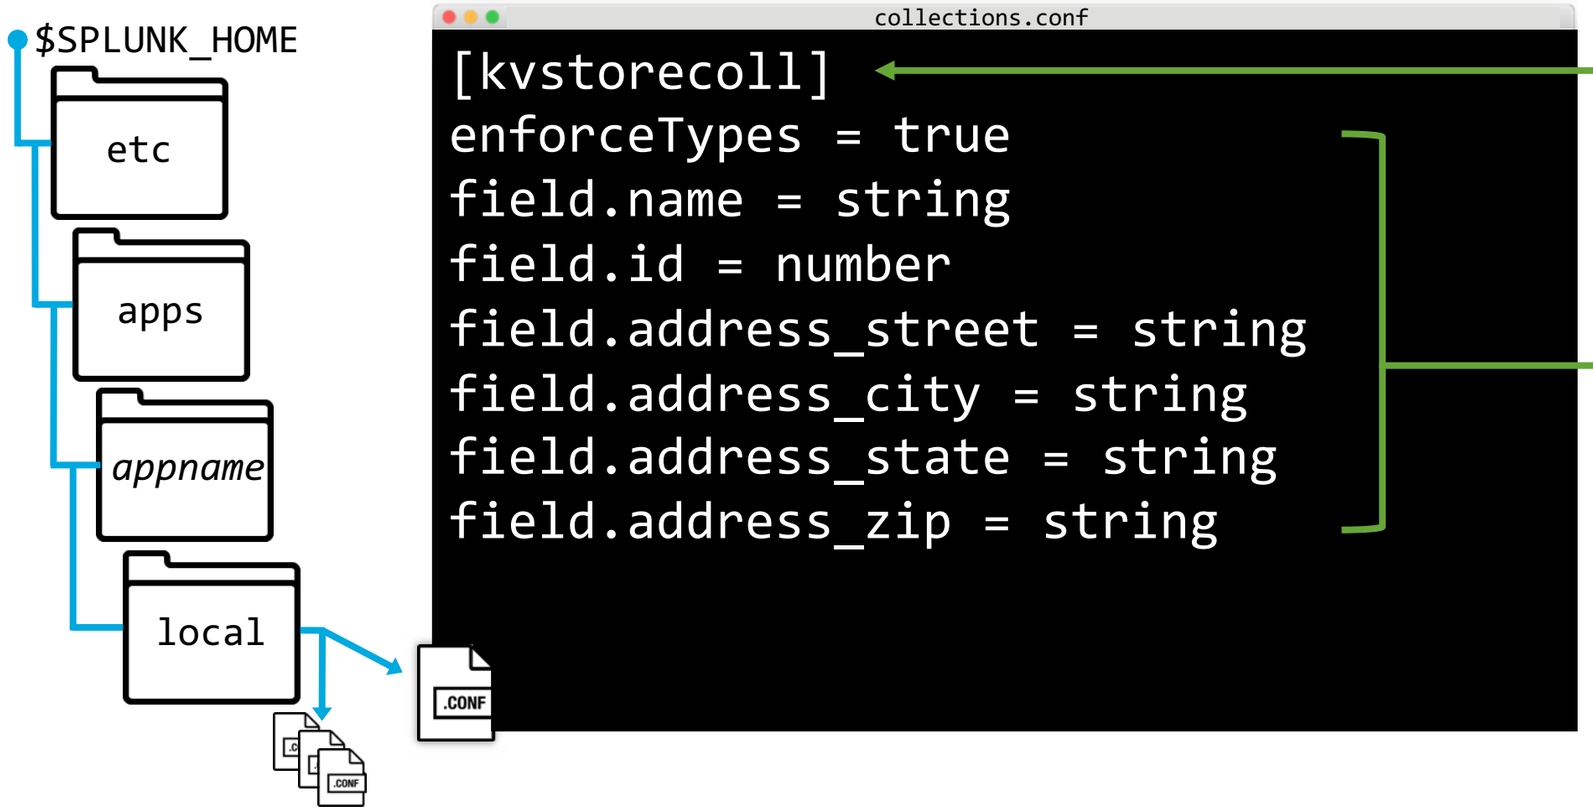

A `[collection_name]` is required

Admins must define the data types for each field and have the option to enforce those data types

Data type options include: **`number`**, **`bool`**, **`string`**, **`time`**

Note

When data type values are enforced, any invalid value added to a collection causes record insertion to fail.

 | 12 August 2022

Set Up a KV Store: `collections.conf` (cont.)

- Enforcing data types is useful if you want to:
  - Guarantee a field is always treated as a specific data type
  - Improve the collection's accelerations (beyond the scope of this course)
- For example, an admin would create the following configuration stanza to enforce the data types of this JSON record

collections.conf

```
[kvstorecoll]
enforceTypes = true
field.name = string
field.id = number
field.address_street = string
field.address_city = string
field.address_state = string
field.address_zip = string
```

```
{
      "name" : "Splunk Seattle",
      "id" : 123,
      "address" :
      {
            "street" : "1730 Minor Avenue",
            "city" : "Seattle",
            "state" : "WA",
            "zip" : "98101"
      }
}
```

Sample **JSON** data

Set Up a KV Store: Create a Definition

Settings > Lookups > Lookup definitions > **New Lookup Definition**

1. Choose Destination app
2. Enter Name that will be used in the search string
3. Change to “KV Store”
4. Enter Collection Name as defined in `collections.conf`
5. List all fields supported by the lookup
6. Save

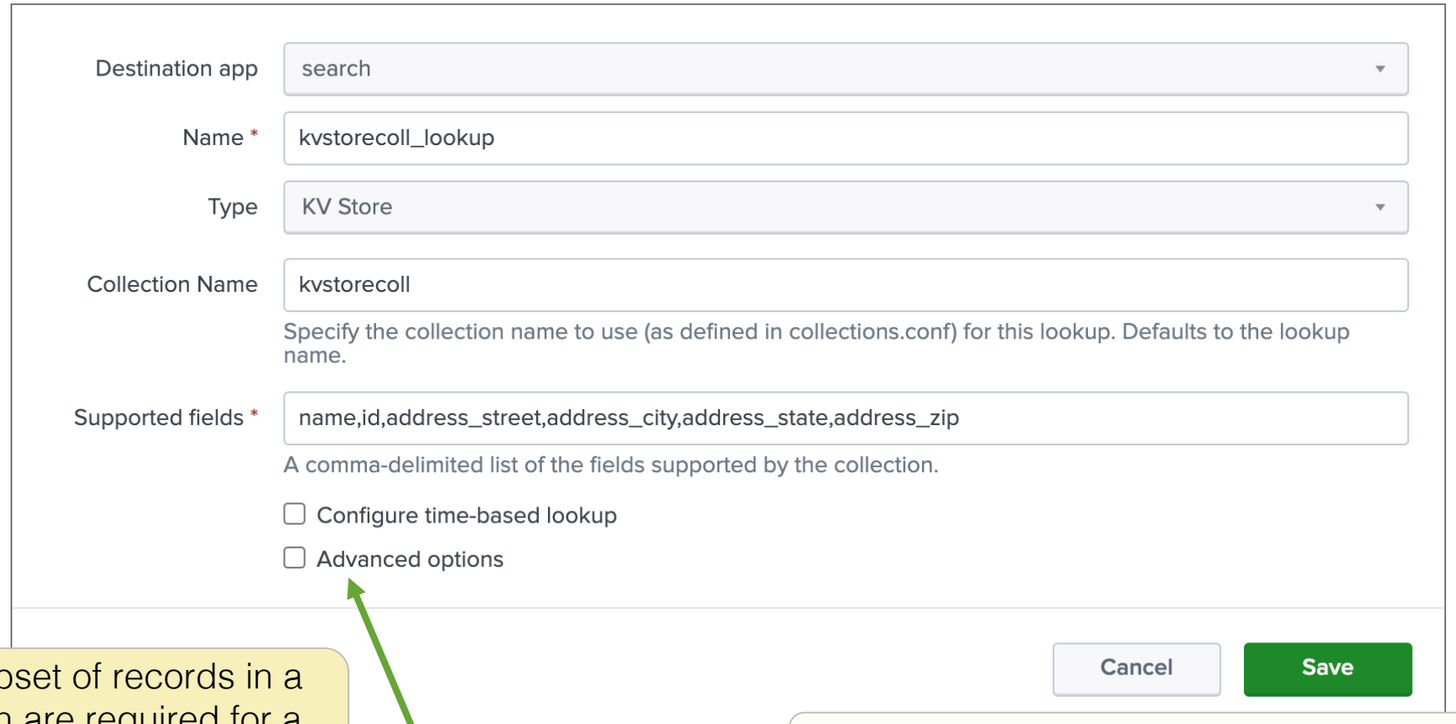

If only results from a subset of records in a large KV store collection are required for a search, improve performance by filtering

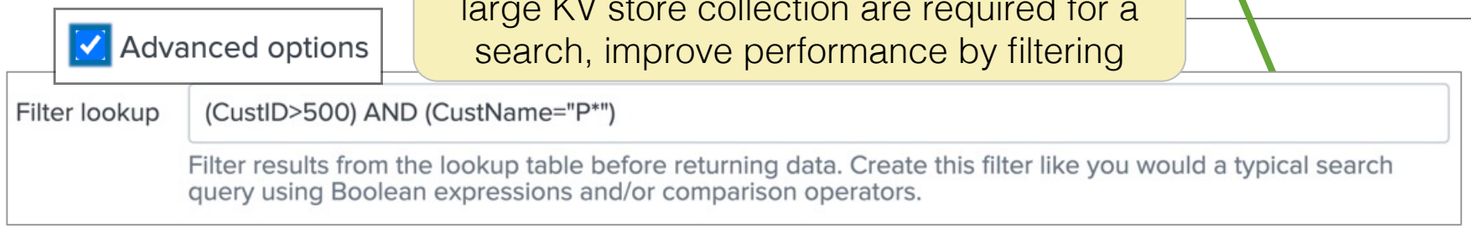

Note

Each collection must have at least two Supported fields. One of these fields must match the values of a field in your event data.

Set Up a KV Store: Populating

Option 1: Use `outputlookup` to write search results into a specific KV Store collection

```
...
|outputlookup kvstorecoll_lookup
```

…sends results to…

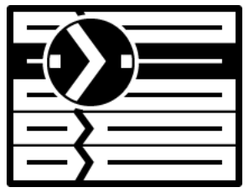

kvstorecoll_lookup

Note

You must have access to the collection in order to write results using the `outputlookup` command.

Option 2: Use Splunk REST API

```
curl -k -u admin:yourpassword \
     https://localhost:8089/servicesNS/nobody/kvst
     oretest/storage/collections/data/kvstorecoll \
     -H 'Content-Type: application/json' \
     -d '{"name": "Splunk HQ", "id": 123, "address": {
"street": "250 Brannan Street", "city": "San Francisco",
"state": "CA", "zip": "94107"}}'
```


 | 12 August 2022

CSV Files vs KV Store

	File-based (CSV)	KV Store
Allows for per-record insertion and editing		✓
Suitable for frequent updating		✓
Allows for data type enforcement		✓
Allows for field accelerations		✓
Provides REST API access to the data collection		✓
Require a full rewrite of the file to edit values	✓	
Supports case-sensitive field lookups	✓	✓
Supports case-insensitive field lookups	✓	
Uploading lookup file is not mandatory		✓
Allows for multiuser access locking		✓
Matches against values in a KV Store		✓
Matches against values in a .csv file	✓	

Best Practices for Lookups

Best Practices for Lookups

splunk> Best Practice

- Order fields in lookup tables so that ‘key’ field is first (leftmost), followed by other values
- After uploading, validate lookup in search by using:

```
| inputlookup <lookup>
```

key	value
…	…
…	…
…	…
…	…

- For commonly used fields, make lookups automatic
- Use gzipped CSV files or KV Store for large lookups
- Keep your lookups fresh and relevant:
  - Do you really need the lookup table to contain a year’s worth of data or is one week enough?
  - Maintain the lookup table and delete older data if not needed

Best Practices for Lookups (cont.)

splunk> Best Practice

Run a Search > Job > Inspect Job

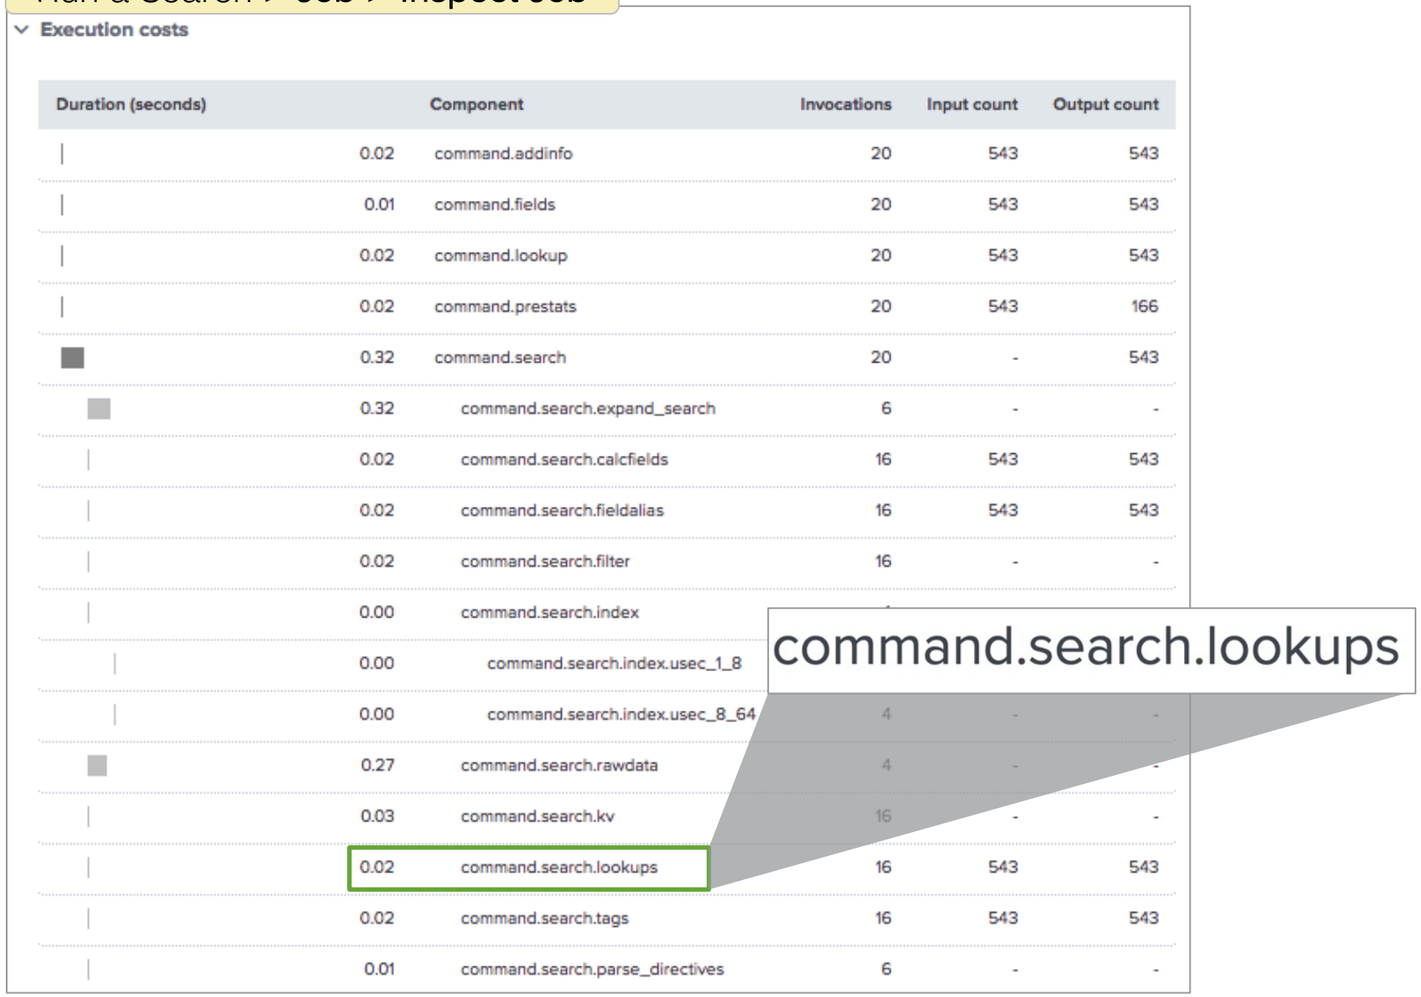

- `command.search.lookups` in job inspector will show how long lookups took to execute
- If there is latency, see if there is one or many lookups being invoked against large files/tables

 | 12 August 2022

Wrap-up Slides

Community

- Splunk Community Portal
  community.splunk.com
  - Answers
  - Discussions
  - Splunk Trust
  - User Groups
  - Ideas
- Splunk Blogs
  splunk.com/blog/
- Splunk Apps
  splunkbase.com
- Splunk Dev Google Group
  groups.google.com/forum/#!forum/splunkdev
- Splunk Docs on Twitter
  twitter.com/splunkdocs
- Splunk Dev on Twitter
  twitter.com/splunkdev
- Splunk Live!
  splunklive.splunk.com
- .conf
  conf.splunk.com

Support Programs

- **Web**
  - Documentation: dev.splunk.com and docs.splunk.com
  - Wiki: wiki.splunk.com
- **Splunk Lantern**
  Guidance from Splunk experts
  - lantern.splunk.com
- **Global Support**
  Support for critical issues, a dedicated resource to manage your account – 24 x 7 x 365
  - Web: splunk.com/index.php/submit_issue
- **Enterprise, Cloud, ITSI, Security Support**
  - Web: splunk.com/en_us/about-splunk/contact-us.html#tabs/customersupport
  - Phone: (855) SPLUNK-S or (855) 775-8657

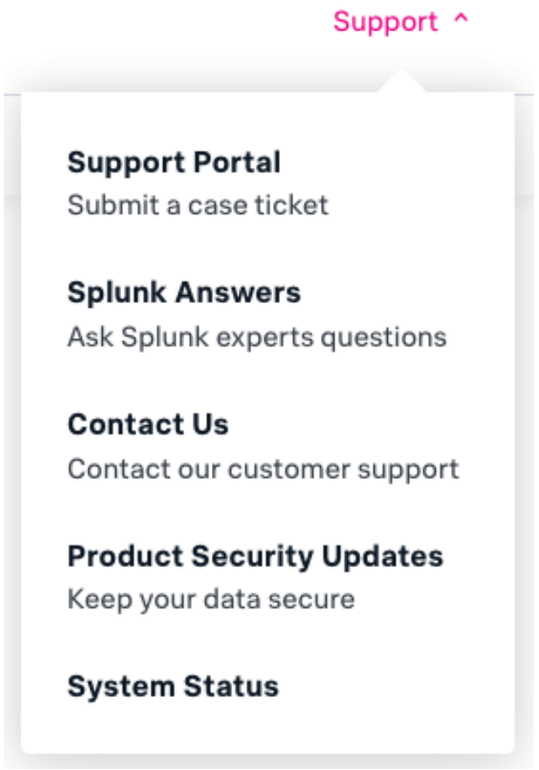

Learning Paths (cont.)

Knowledge Manager - Recommended Courses

Free eLearning courses are in blue and courses with an * are present in both learning paths.

- What is Splunk *
- Introduction to Splunk *
- Using Fields *
- Introduction to Knowledge Objects
- Creating Knowledge Objects
- Creating Field Extractions
- Enriching Data with Lookups
- Data Models
- Introduction to Dashboards
- Dynamic Dashboards
- Using Choropleth
- Search Optimization *

Learning Paths

Search Expert - Recommended Courses

Free eLearning courses are in blue and courses with an * are present in both learning paths.

- What is Splunk *
- Introduction to Splunk *
- Using Fields *
- Scheduling Reports and Alerts
- Visualizations
- Statistical Processing
- Working with Time
- Comparing Values
- Result Modification
- Leveraging Lookups and Subsearches
- Correlation Analysis
- Search Under the Hood
- Multivalue Fields
- Search Optimization *

Splunk Certification

Offerings & Requirements

Splunk Core and Beyond

Regardless of which Splunk product you use, it all starts with Splunk Core

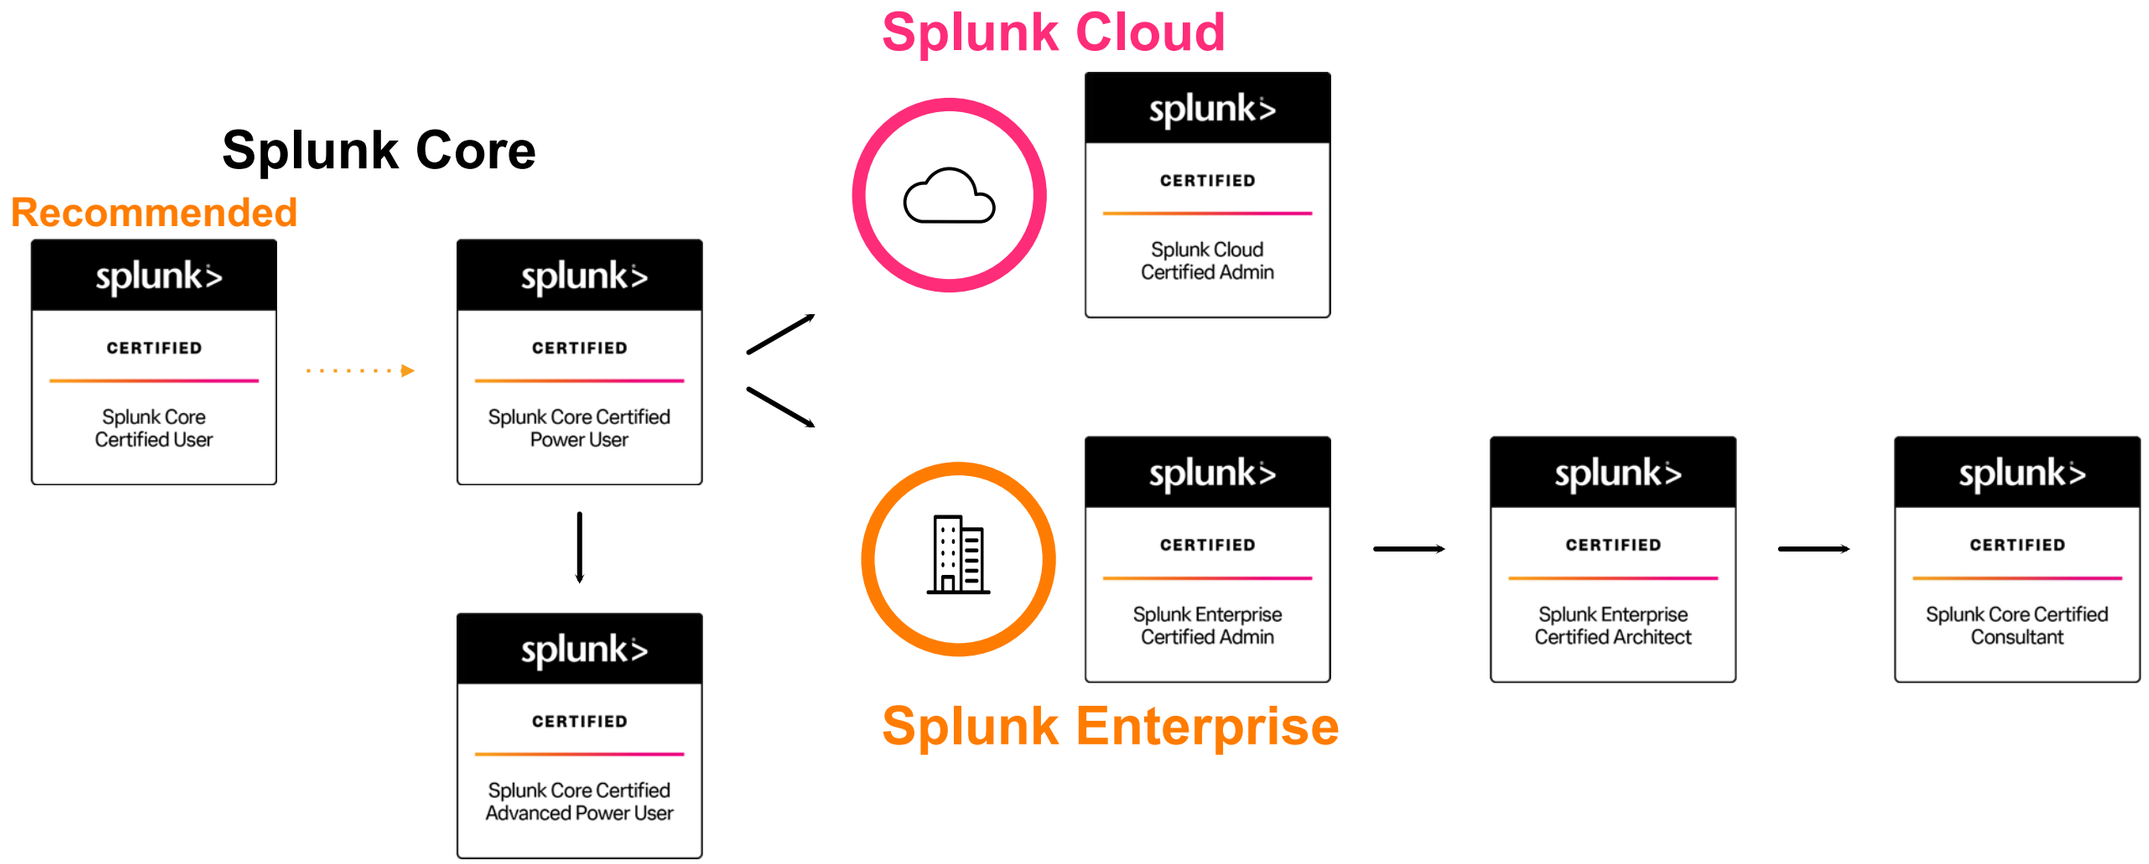

Splunk Core Certified User

This entry-level certification demonstrates an individual's basic ability to navigate and use Splunk software

Prerequisite Certification(s):

- None

Prerequisite Course(s):

- None

Splunk Core Certified User Exam

Time to study! We suggest candidates looking to prepare for this exam complete Fundamentals 1 **or** the following courses:

- What is Splunk?
- Intro to Splunk
- Using Fields
- Scheduling Reports and Alerts
- Visualizations
- Statistical Processing
- Working with Time
- Leveraging Lookups and Subsearches
- Search Optimization
- Enriching Data with Lookups
- Data Models

See here for registration assistance.

Congratulations! You are a...

Recommended Next Step

- Splunk Core Certified Power User

 | 12 August 2022

Splunk Core Certified Advanced Power User

This certification demonstrates an individual's ability to generate complex searches, reports, and dashboards with Splunk's core software to get the most out of their data

Prerequisite Certification(s):

- Splunk Core Certified Power User

Prerequisite Course(s):

- None

Splunk Core Certified Advanced Power User Exam

Time to study! We suggest candidates looking to prepare for this exam complete Fundamentals 3, Creating Dashboards, and Advanced Searching & Reporting **or** the following courses:

- Using Fields
- Working with Time
- Comparing Values
- Result Modification
- Leveraging Lookups and Subsearches
- Correlation Analysis
- Search Under the Hood
- Multivalue Fields
- Search Optimization
- Creating Field Extractions
- Enriching Data with Lookups
- Data Models
- Using Choropleth
- Introduction to Dashboards
- Dynamic Dashboards

See here for registration assistance.

Congratulations! You are a...

Recommended Next Steps

- Splunk Enterprise Certified Admin
- Splunk Cloud Certified Admin

 | 12 August 2022

Thank You

 | 12 August 2022